

Final Notebook (Including Cross-Validation) (70-15-15)

April 13, 2024

1 Ovarian Cancer Modelling - A comparative Analysis of Baseline and Ensemble Methods

1.1 Introduction

This notebook is designed to develop machine-learning models for detecting ovarian cancer. The main aim of this research is to utilize various machine-learning algorithms for detecting **Ovarian Cancer**. The code is organized into several sections.

1. **The Data Description section** describes the dataset used in this research.
2. The Data Wrangling section preprocesses the dataset to transform raw data into a clean dataset suitable for training the machine learning model.
3. The Exploratory Data Analysis section visualizes key features in the provided dataset to gain an insight into the data graphically.
4. The Preliminary Data Analysis section explores the provided dataset and performs basic statistical analysis to understand the data better.
5. The Feature Engineering section extracts relevant features from the data to improve the model's accuracy.
6. The Model Training section trains the machine learning model using several ensemble learning algorithms and evaluates their performance based on various metrics.
7. The Analysis of Ensemble Methods section analyzes the factors that contribute to the performance of the selected baseline algorithms.
8. The Analysis of Ensemble Methods section analyzes the factors that contribute to the performance of the selected ensemble learning algorithms.
9. Assessing how the Number of Features Impacts Model Performance Section systematically analyzes the factors contributing to the performance of selected Decision Tree models using varying numbers of features determined by the MRMR method.

1.1.1 Data Description

This section of the code provides sufficient information about the dataset used for this research.

The Third Affiliated Hospital of Soochow University provided the dataset for the study, which includes 349 individuals. The data were collected between July 2011 and July 2018, and they were divided into two groups: 178 patients with benign ovarian tumors and 171 patients with

ovarian cancer ([Kaggle](#), accessed on 15 January 2024). 49 features in all, derived through pathology diagnosis, were included in the dataset. These 49 predictor factors included information on age and menopause, as well as 22 basic chemical tests, 19 normal blood tests, and 6 tumor markers. Prior to surgery, none of the patients had received chemotherapy or radiotherapy, and all underwent postoperative case diagnosis. Using standards from the World Health Organization, the histological diagnosis was categorized.

Biomarker	Biomarker Name
MPV	Mean platelet volume
BASO#	Basophil Cell Count
PHOS	phosphorus
GLU.	glucose
CA72-4	Carbohydrate antigen 72-4
K	kalium
AST	Aspartate aminotransferase
BASO%	Basophil Cell ratio
Mg	magnesium
CL	chlorine
CEA	Carcinoembryonic antigen
EO#	eosinophil count
CA19-9	Carbohydrate antigen 19-9
ALB	albumin
IBIL	Indirect bilirubin
GGT	Gama glutamyltransferasey
MCH	Mean corpuscular hemoglobin
GLO	globulin
DBIL	direct bilirubin
RDW	red blood cell distribution width
PDW	Platelet distribution width
CREA	creatinine
AFP	alpha-fetoprotein
HGB	hemoglobin
Na	Natrium
HE4	human epididymis protein 4
LYM#	lymphocyte count
CA125	Carbohydrate antigen 125
BUN	blood urea nitrogen
LYM%	lymphocyte ratio
Ca	calcium
AG	Anion gap
MONO#	mononuclear cell count
PLT	platelet count
NEU	neutrophil ratio
EO%	eosinophil ratio
TP	Total protein
UA	urie acid
RBC	Red blood cell count

Biomarker	Biomarker Name
PCT	thrombocytocrit
CO2CP	carban dioxide-combining Power
TBIL	total bilirubin
HCT	hematocrit
MONO%	monocyte ratio
MCV	mean corpuscular volume
ALP	Alkaline phosphatase

```
[1]: import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import seaborn as sns # data visualization
import matplotlib.pyplot as plt # data visualization
import shutil

import joblib
import pickle
import os
import time
import warnings
warnings.filterwarnings('ignore')

import mlflow

mlflow.set_tracking_uri(uri="http://127.0.0.1:5000")
mlflow.autolog()

from mlflow.models import infer_signature

import plotly.express as px #data visualization
import plotly.io as pio # plot rendering
pio.renderers.default = 'jupyterlab'
#pio.renderers.default = "plotly_mimetype+notebook"
import plotly.figure_factory as ff # data visualization
import plotly.graph_objects as go # data visualization
from plotly.subplots import make_subplots # data visualization

from mrmr import mrmr_classif # Importing mrmr lib for top features selection.

from IPython.display import display, Javascript # utilizing JavaScript for
rendering charts
```

2024/04/13 12:32:05 INFO mlflow.tracking.fluent: Autologging successfully enabled for statsmodels.

2024/04/13 12:32:05 WARNING mlflow.utils.autologging_utils: You are using an unsupported version of sklearn. If you encounter errors during autologging, try upgrading / downgrading sklearn to a supported version, or try upgrading MLflow.

2024/04/13 12:32:06 INFO mlflow.tracking.fluent: Autologging successfully enabled for sklearn.

```
[2]: # Create a new MLflow Experiment
```

```
mlflow.set_experiment("Ovarian Cancer Prediction - (70-15-15)")
```

2024/04/13 12:32:06 INFO mlflow.tracking.fluent: Experiment with name 'Ovarian Cancer Prediction - (70-15-15)' does not exist. Creating a new experiment.

```
[2]: <Experiment: artifact_location='mlflow-artifacts:/164884331870191523',
creation_time=1713007926481, experiment_id='164884331870191523',
last_update_time=1713007926481, lifecycle_stage='active', name='Ovarian Cancer
Prediction - (70-15-15)', tags={}>
```

```
[3]: # Pandas display options for easy viewing of data frames
```

```
pd.set_option('display.width', 150)
```

```
[4]: # Reading the dataset
```

```
df_cancer = pd.read_csv(
    "https://docs.google.com/spreadsheets/d/e/
    ↪2PACX-1vRveqCgY-ndsv7PmJwBSb5sN3ZmmbQ6kKFK79q8H7powQoj1h4BzRx5qHzzSdIRDQ/pub?
    ↪gid=1679652868&single=true&output=csv",
    dtype=str
)
# Remove trailing whitespace from all string columns
df_cancer = df_cancer.apply(lambda x: x.str.rstrip() if x.dtype == "object"
    ↪else x)
```

```
[5]: # Printing the first 5 rows of the data
```

```
df_cancer.head()
```

```
[5]:  SUBJECT_ID  AFP      AG Age  ALB ALP ALT AST BASO# BASO%  ...  NEU  PCT
PDW  PHOS  PLT   RBC   RDW  TBIL   TP   UA
0      1   3.58  19.36  47  45.4  56  11  24  0.01   0.3  ...  76.2  0.09
13.4  1.46  74  2.64  13.7   5.5  73.9  396.4
1      2  34.24  23.98  61  39.9  95   9  13  0.02   0.3  ...  76.5   0.3
11.2  1.09  304  4.89  12.7   6.8   72  119.2
2      3   1.50  18.4  39  45.4  77   9  18  0.03   0.6  ...  69.7  0.13
15.2  0.97  112  4.62   12  14.8  77.9  209.2
3      4   2.75  16.6  45  39.2  26  16  17  0.05  0.74  ...  65.5  0.25
17.4  1.25  339  4.01  14.6  10.9  66.1  215.6
4      5   2.36  19.97  45   35  47  21  27  0.01   0.1  ...  59.5  0.28
11.9  0.94  272  4.4  13.4   5.3  66.5   206
```

[5 rows x 51 columns]

```
[6]: # Printing the summary statistics of each feature/column in the data
```

```
df_cancer.describe().T
```

```

[6]:
      count unique    top freq
SUBJECT_ID    349    349    414    1
AFP           327    237    0.61    5
AG            348    307   18.58    3
Age           349     62     45   11
ALB           339    172   42.6     8
ALP           339     97     71   11
ALT           339     46     16   30
AST           339     42     13   31
BASO#         349     12    0.02   75
BASO%         349     75     0.2   48
BUN           349    246     3.8     5
Ca            349     84     2.5   13
CA125         332    326   1319     2
CA19-9        325    300  <0.600     6
CA72-4        109    100     0.2     5
CEA           327    207     1.11    5
CL            349    117    99.3   11
CO2CP         348    104    24.6   12
CREA          349    146     56   14
TYPE          349      2      1  178
DBIL          339     60     2.5   23
EO#           349     32      0   39
EO%           349    108      0   23
GGT           339     58     12   28
GLO           339    150    32.1     8
GLU.          349    203     4.5     5
HCT           349    146    0.386     8
HE4           329    321   219.1     2
HGB           349     92    123   17
IBIL          339    100     5.4   12
K             349    131     4.3   10
LYM#          349    172     1.69    6
LYM%          349    243    30.4     5
MCH           349    107    30.5   14
MCV           349    162    91.5     7
Menopause     349      2      0  230
Mg            349     61      1   22
MONO#         349     68    0.36   17
MONO%         349    143     5.1   16
MPV           347    135    10.9   16
Na            349    107   138.6     8
NEU           258    202    65.5     4
PCT           347    113     0.23   22
PDW           347    121    17.4     8
PHOS          349     82     1.2   10
PLT           349    195    247     6

```

RBC	349	148	4.24	8
RDW	349	84	13.2	18
TBIL	339	125	5.7	8
TP	339	182	77	7
UA	349	326	229.2	3

```
[7]: # Printing the shape of the data
df_cancer.shape
```

```
[7]: (349, 51)
```

```
[8]: # Printing the concise summary of the data
df_cancer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 349 entries, 0 to 348
Data columns (total 51 columns):
#   Column          Non-Null Count  Dtype
---  -
0   SUBJECT_ID      349 non-null    object
1   AFP              327 non-null    object
2   AG              348 non-null    object
3   Age             349 non-null    object
4   ALB             339 non-null    object
5   ALP             339 non-null    object
6   ALT             339 non-null    object
7   AST             339 non-null    object
8   BASO#           349 non-null    object
9   BASO%           349 non-null    object
10  BUN             349 non-null    object
11  Ca              349 non-null    object
12  CA125           332 non-null    object
13  CA19-9          325 non-null    object
14  CA72-4          109 non-null    object
15  CEA             327 non-null    object
16  CL              349 non-null    object
17  CO2CP           348 non-null    object
18  CREA            349 non-null    object
19  TYPE            349 non-null    object
20  DBIL            339 non-null    object
21  EO#             349 non-null    object
22  EO%             349 non-null    object
23  GGT             339 non-null    object
24  GLO             339 non-null    object
25  GLU.            349 non-null    object
26  HCT             349 non-null    object
27  HE4             329 non-null    object
28  HGB             349 non-null    object
```

29	IBIL	339 non-null	object
30	K	349 non-null	object
31	LYM#	349 non-null	object
32	LYM%	349 non-null	object
33	MCH	349 non-null	object
34	MCV	349 non-null	object
35	Menopause	349 non-null	object
36	Mg	349 non-null	object
37	MONO#	349 non-null	object
38	MONO%	349 non-null	object
39	MPV	347 non-null	object
40	Na	349 non-null	object
41	NEU	258 non-null	object
42	PCT	347 non-null	object
43	PDW	347 non-null	object
44	PHOS	349 non-null	object
45	PLT	349 non-null	object
46	RBC	349 non-null	object
47	RDW	349 non-null	object
48	TBIL	339 non-null	object
49	TP	339 non-null	object
50	UA	349 non-null	object

dtypes: object(51)

memory usage: 139.2+ KB

1.1.2 Data Wrangling

This section of the code is responsible for preparing the dataset for analysis by cleaning, transforming, and restructuring the data into a usable format. This section involves handling missing data, dealing with outliers, and transforming variables to ensure they meet the assumptions of the analysis method. The goal is to create a reliable dataset that maximizes accuracy when using machine learning algorithms. Data wrangling is a critical step in the data analysis process, as the accuracy of the results depends heavily on the quality of the dataset used.

```
[9]: # Before converting column types, the column values should be handled for data
      ↪inconsistency
```

```
df_cancer.loc[df_cancer['AFP'] == '>1210.00', 'AFP'] = '1210.00'
df_cancer.loc[df_cancer['AFP'] == '>1210', 'AFP'] = '1210.00'
df_cancer.loc[df_cancer['CA125'] == '>5000.00', 'CA125'] = '5000.00'
df_cancer.loc[df_cancer['CA19-9'] == '>1000.00', 'CA19-9'] = '1000.00'
df_cancer.loc[df_cancer['CA19-9'] == '>1000', 'CA19-9'] = '1000.00'
df_cancer.loc[df_cancer['CA19-9'] == '<0.600', 'CA19-9'] = '0.5'
```

```
[10]: # Convert object columns to float columns
for col in (
    df_cancer.drop(["TYPE", "SUBJECT_ID",], axis=1)
    .select_dtypes(include=["object"])
    .columns
```

```

):
    df_cancer[col] = df_cancer[col].astype('float')

# Convert target column to integer
df_cancer['TYPE'] = df_cancer['TYPE'].astype('int64')

```

```

[11]: # Computing the ratio of missing data in each column
missing_ratio = df_cancer.isnull().sum()

# Displaying the ratio of missing data in each column
missing_ratio

```

```

[11]: SUBJECT_ID      0
AFP                  22
AG                   1
Age                  0
ALB                  10
ALP                  10
ALT                  10
AST                  10
BASO#                 0
BASO%                 0
BUN                   0
Ca                    0
CA125                 17
CA19-9                24
CA72-4               240
CEA                   22
CL                     0
CO2CP                 1
CREA                  0
TYPE                  0
DBIL                  10
EO#                   0
EO%                   0
GGT                   10
GLO                   10
GLU.                  0
HCT                   0
HE4                   20
HGB                   0
IBIL                  10
K                     0
LYM#                  0
LYM%                  0
MCH                   0
MCV                   0

```


Menopause	0
Mg	0
MONO#	0
MONO%	0
MPV	2
Na	0
NEU	91
PCT	2
PDW	2
PHOS	0
PLT	0
RBC	0
RDW	0
TBIL	10
TP	10
UA	0

dtype: int64

```
[12]: # Computing the ratio of missing data in each column
missing_ratio = df_cancer.isnull().mean()

# Displaying the ratio of missing data in each column
missing_ratio
```

```
[12]: SUBJECT_ID    0.000000
AFP              0.063037
AG              0.002865
Age             0.000000
ALB             0.028653
ALP             0.028653
ALT             0.028653
AST             0.028653
BASO#           0.000000
BASO%           0.000000
BUN             0.000000
Ca              0.000000
CA125           0.048711
CA19-9          0.068768
CA72-4          0.687679
CEA             0.063037
CL              0.000000
CO2CP           0.002865
CREA            0.000000
TYPE            0.000000
DBIL            0.028653
EO#             0.000000
EO%             0.000000
```

GGT	0.028653
GLO	0.028653
GLU.	0.000000
HCT	0.000000
HE4	0.057307
HGB	0.000000
IBIL	0.028653
K	0.000000
LYM#	0.000000
LYM%	0.000000
MCH	0.000000
MCV	0.000000
Menopause	0.000000
Mg	0.000000
MONO#	0.000000
MONO%	0.000000
MPV	0.005731
Na	0.000000
NEU	0.260745
PCT	0.005731
PDW	0.005731
PHOS	0.000000
PLT	0.000000
RBC	0.000000
RDW	0.000000
TBIL	0.028653
TP	0.028653
UA	0.000000

dtype: float64

```
[13]: # Before handling missing data, let's put the missing data in another variable
      ↪to perform analysis later
df_cancer_missing = df_cancer.copy()

# Dropping columns with a missing data ratio greater than 0.5
cols_to_drop = ["CA72-4", "CA19-9", "AFP", "CEA", "HE4", "SUBJECT_ID"]
df_cancer = df_cancer.drop(cols_to_drop, axis=1)

# get columns with missing data
cols_with_missing = [col for col in df_cancer.columns if df_cancer[col].
    ↪isnull().any()]

# impute missing data with the median value
for col in cols_with_missing:
    median_val = df_cancer[col].median()
    df_cancer[col].fillna(median_val, inplace=True)
```

```
# Displaying the updated missing data ratio  
df_cancer.isnull().mean()
```

```
[13]: AG          0.0  
      Age          0.0  
      ALB          0.0  
      ALP          0.0  
      ALT          0.0  
      AST          0.0  
      BASO#        0.0  
      BASO%        0.0  
      BUN          0.0  
      Ca           0.0  
      CA125        0.0  
      CL           0.0  
      CO2CP        0.0  
      CREA         0.0  
      TYPE         0.0  
      DBIL         0.0  
      EO#          0.0  
      EO%          0.0  
      GGT          0.0  
      GLO          0.0  
      GLU.         0.0  
      HCT          0.0  
      HGB          0.0  
      IBIL         0.0  
      K            0.0  
      LYM#         0.0  
      LYM%         0.0  
      MCH          0.0  
      MCV          0.0  
      Menopause    0.0  
      Mg           0.0  
      MONO#        0.0  
      MONO%        0.0  
      MPV          0.0  
      Na           0.0  
      NEU          0.0  
      PCT          0.0  
      PDW          0.0  
      PHOS         0.0  
      PLT          0.0  
      RBC          0.0  
      RDW          0.0  
      TBIL         0.0  
      TP           0.0
```

```
UA
dtype: float64
```


```
[14]: # Calculate the ratio of missing values in each column
missing_ratio = df_cancer_missing.isnull().sum() / len(df_cancer_missing)

# Filter columns with missing values
missing_ratio = missing_ratio[missing_ratio > 0]

# Create the bar chart using Plotly
fig = px.bar(x=missing_ratio.index, y=missing_ratio.values, text_auto=True,
             color=missing_ratio.index)

# Update the layout of the plot
fig.update_layout(
    title="Ratio of missing values in columns with missing values",
    xaxis_title="Column name",
    yaxis_title="Ratio of missing values",
)

# Show the plot
fig.show()
```



Of all the feature variables, Carbohydrate antigen 72-4 (CA72-4) had ~69% missing observations, Neutrophil ratio (NEU) had 26%, and Carbohydrate antigen 19-9 (~7%) being the top 3 features with the highest percentage of missing observations.

1.1.3 Feature Engineering

This section of the code refers to the process of selecting and transforming the relevant features of the data to create new features that better represent the problem domain. The purpose of feature engineering is to improve the performance of machine learning algorithms by reducing the noise in the data, increasing the accuracy of the predictions, and making the model more interpretable. It requires a deep understanding of the problem domain and the data being used, as well as knowledge of the available feature engineering techniques and their impact on the model's performance.

```
[15]: # Map the 'TYPE' column to "0 - Benign Ovarian Tumors" and '1 - Ovarian Cancer'
df_cancer["Type_Label"] = df_cancer["TYPE"].map(
    {0: "Ovarian Cancer", 1: "Benign Ovarian Tumors"}
)
# Map the 'Menopause' column to "0 - No" and '1 - Yes'
df_cancer['Menopause_Label'] = df_cancer['Menopause'].map({0: 'No', 1: 'Yes'})
df_cancer.head()
```

```
[15]:      AG  Age  ALB  ALP  ALT  AST  BASO#  BASO%  BUN  Ca  ...  PDW
PHOS  PLT  RBC  RDW  TBIL  TP    UA      Type_Label  \
0  19.36  47.0  45.4  56.0  11.0  24.0  0.01  0.30  5.35  2.48  ...  13.4
1.46  74.0  2.64  13.7  5.5  73.9  396.4  Ovarian Cancer
1  23.98  61.0  39.9  95.0  9.0  13.0  0.02  0.30  3.21  2.62  ...  11.2
1.09  304.0  4.89  12.7  6.8  72.0  119.2  Ovarian Cancer
2  18.40  39.0  45.4  77.0  9.0  18.0  0.03  0.60  3.80  2.57  ...  15.2
0.97  112.0  4.62  12.0  14.8  77.9  209.2  Ovarian Cancer
3  16.60  45.0  39.2  26.0  16.0  17.0  0.05  0.74  5.27  2.35  ...  17.4
1.25  339.0  4.01  14.6  10.9  66.1  215.6  Ovarian Cancer
4  19.97  45.0  35.0  47.0  21.0  27.0  0.01  0.10  4.89  2.48  ...  11.9
0.94  272.0  4.40  13.4  5.3  66.5  206.0  Ovarian Cancer

Menopause_Label
0          No
1          Yes
2          No
3          Yes
4          No
```

[5 rows x 47 columns]

```
[16]: # split data into features (X) and target (y)
X = df_cancer.drop(['TYPE', 'Type_Label', 'Menopause_Label'], axis=1)
y = df_cancer['TYPE']
```

```
[17]: # select top 20 features using MRMR
top_features = mrmr_classif(X=X, y=y, K=20)
```

100%|

| 20/20 [00:00<00:00, 51.20it/s]

```
[18]: top_features
```

```
[18]: ['Age',
      'IBIL',
      'NEU',
      'Menopause',
      'ALB',
```

```
'CA125',
'GLO',
'LYM%',
'AST',
'HGB',
'PLT',
'ALP',
'LYM#',
'PCT',
'Ca',
'MONO#',
'TBIL',
'GLU.',
'MCH',
'Na']
```

The top 20 features contributing to the target variable (TYPE) are selected using the mRMR algorithm. The selected features are: 1. Age 2. Indirect Bilirubin (IBIL) 3. Neutrophil ratio (NEU) 4. Menopause 5. Albumin (ALB) 6. Carbohydrate antigen 125 (CA125) 7. Globulin (GLO) 8. Lymphocyte ratio (LYM%) 9. Aspartate aminotransferase (AST) 10. Hemoglobin (HGB) 11. Platelet count (PLT) 12. Alkaline phosphate (ALP) 13. Lymphocyte count (LYM#) 14. Thrombocytocrit (PCT) 15. Calcium (ca) 16. Mononuclear cell count (MONO#) 17. Total Bilirubin (TBIL) 18. Glucose (GLU) 19. Mean corpuscular hemoglobin (MCH) 20. Natrium (Na)

1.1.4 Preliminary Data Analysis

This section of the code involves an initial examination of the dataset to understand its structure, contents, and quality. This includes checking for the extent of missing or erroneous data, exploring the distribution of the variables, identifying any outliers, and computing summary statistics. The purpose of this section is to gain insights into the dataset and inform subsequent data processing steps. It also involves visualizing the data using various plotting techniques to reveal patterns or relationships between the variables.

```
[19]: # Visualizing the correlation between the features in the training data
selected_cols = top_features + ["TYPE"]

# Selecting the row or column corresponding to the target column
target_corr = df_cancer[selected_cols].corr()['TYPE']
target_corr = np.sqrt(target_corr ** 2) # find the magnitude of each
    ↳ correlation, at this instance, we are only interested in the magnitude
    ↳ of each correlation, and not its direction.

# Displaying the correlation values of each feature with the target column in
    ↳ descending order
print("Correlation of selected features with target column (sorted in
    ↳ descending order) \n")
target_corr.sort_values(ascending=False)
```

Correlation of selected features with target column (sorted in descending order)

```
[19]: TYPE          1.000000
      Age           0.514098
      Menopause     0.455770
      ALB           0.375415
      CA125         0.372262
      NEU           0.353062
      LYM%          0.315035
      PLT           0.270182
      LYM#          0.256494
      PCT           0.243719
      AST           0.215888
      ALP           0.213249
      MONO#         0.200536
      IBIL          0.200451
      HGB           0.197863
      TBIL          0.195921
      GLO           0.195630
      Ca            0.187119
      GLU.          0.179048
      MCH           0.166818
      Na            0.143849
      Name: TYPE, dtype: float64
```

```
[20]: # Create a correlation matrix
corr_matrix = np.around(np.sqrt(df_cancer[selected_cols].corr()**2), 2) # Use
    ↳ the correlation magnitude instead of the direction

# Create the heatmap using Plotly
fig = ff.create_annotated_heatmap(
    z=corr_matrix.to_numpy(),
    x=corr_matrix.columns.tolist(),
    y=corr_matrix.columns.tolist(),
    annotation_text=corr_matrix.to_numpy().astype(str),
    colorscale="RdBu_r", # Using a similar colorscale to 'coolwarm'
    showscale=True,
)

# Update the layout of the plot
fig.update_layout(
    title="Correlation Heatmap of Top Features",
    width=1200,
    height=1000,
    xaxis=dict(tickangle=-45),
    yaxis=dict(tickmode="array", tickvals=np.arange(len(selected_cols))),
```


```

margin=dict(l=200, r=200, t=100, b=100),
)

# Show the plot
fig.show()

```

Correlation Heatmap of Top Features



Based on the available data, and the correlation result above, it is safe to say that women are at a higher risk of developing ovarian cancer as they age. The top three factors are Age (51%), Menopause (46%), and Albumin (38%) respectively, with the least three factors being Calcium (19%), Carcinoembryonic antigen (17%), and Carbohydrate antigen 19-9 (15%).

```

[21]: fig = px.scatter_matrix(
        df_cancer[top_features],
        dimensions=top_features,
        width=1200,
        height=1500


```




```

)
fig.update_layout(
    title='Scatter Plot Matrix of Selected Features'
)
fig.show()

```




```
[22]: # Visualizing the distribution of the target variable in the training data
fig = px.histogram(df_cancer, x='Type_Label', text_auto=True, orientation='v',
    color='Type_Label')
fig.update_layout(
    xaxis_title='Cancer Type',
    yaxis_title='Number of Patients',
    title='Distribution of Target Variable'
)
fig.show()
```




From the above, it can be seen that there is a class balance between the number of patients with Ovarian cancer (0), and patients with Benign Ovarian Tumors (1).

```
[23]: # Plotting the distribution of cancer types by menopause status
fig = px.histogram(df_cancer, x='Menopause_Label', color='Type_Label',
    barmode='group', text_auto=True)
fig.update_layout(
    title='Distribution of Cancer Types by Menopause Status',
    yaxis_title = "Number of Patients",
    xaxis_title = "Menopause"
)
fig.show()
```



Women at the menopause stage have higher records of Ovarian Cancer than women not at this stage. From the above chart, out of 171 patients with Ovarian cancer, **56.14% (96 patients)** have Ovarian cancer and are at the menopause stage. In contrast, out of the 178 patients with Benign Ovarian Tumors, only **12.92% (23 patients)** are at the menopause stage.

```
[24]: # Plotting the distribution of cancer types by menopause status
fig = px.histogram(df_cancer, x='Age', color='Type_Label', barmode='group',
    ↪text_auto=True)
fig.update_layout(
    title="Distribution of Patients' Age by Cancer Type",
    yaxis_title = "Number of Patients",
    xaxis_title = "Cancer Type"
)
fig.show()
```




```
[25]: # Filter the data to only include positive values of the "TYPE" target column
positive_data = df_cancer[df_cancer['Type_Label'] == "Ovarian Cancer"]

fig = px.histogram(positive_data, x="Age", text_auto=True)
fig.update_layout(
    xaxis_title="Age",
```

```
yaxis_title="Number of Patients",
title="Distribution of Positive Patients' Age")

fig.show()
```



Women are at more risk of ovarian cancer as they age. From the above chart, women between the age-range of 45 - 64 have more tendencies of developing ovarian cancer with 50 - 54 being the most prevalent age-range, followed by the 60 - 64 age-range.

1.1.5 Model Training

This section of the code involves using machine learning algorithms to build predictive models for detecting ovarian cancer. This section includes selecting appropriate algorithms and splitting the data into training and testing sets, training the models on the training data, and evaluating their performance on the testing data.

The goal of this section is to identify the most accurate and effective ensemble learning method(s) for detecting ovarian cancer. Ensemble learning is a machine learning technique that involves combining multiple models (called baseline models) to improve the overall performance and accuracy of predictions.

Base Models Checked: 1. Logistic Regression

2. SVM

3. KNN

4. Decision Trees

Ensemble Learning Techniques Checked: 1. Voting

2. Stacking

3. Bagging

4. Boosting (XGBoost and GBM)

5. Stacking of Various Ensemble Learning Techniques

```
[26]: from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (
    GradientBoostingClassifier,
    VotingClassifier,
    BaggingClassifier,
    StackingClassifier
)

from sklearn.svm import SVC
from xgboost import XGBClassifier
```

2024/04/13 12:32:14 INFO mlflow.tracking.fluent: Autologging successfully enabled for xgboost.

```
[27]: from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    accuracy_score,
    roc_auc_score,
    roc_curve,
    precision_score,
    auc,
    precision_recall_curve,
    f1_score,
    average_precision_score,
)
```

```
[28]: # Initial split into training (70%) and a temporary set (30%)
X_train, X_temp, y_train, y_temp = train_test_split(
    X, y, test_size=0.3, random_state=42
)

# Split the temporary set into testing (50%) and cross-validation (50%) of the
↳ remaining 30%,
# which translates to 15% of the total data each for test and validation sets.
X_test, X_cv, y_test, y_cv = train_test_split(
    X_temp, y_temp, test_size=0.5, random_state=42
)

# Select top features after splitting
X_train = X_train[top_features]
X_test = X_test[top_features]
X_cv = X_cv[top_features]
```

1.1.6 Analysis of Baseline Algorithms

This section of the code involves evaluating the performance of the baseline methods used in the model training section, by using appropriate evaluation metrics and comparing the results. The purpose is to determine the most effective and efficient method for detecting ovarian cancer tumors and to identify the factors that contribute to the superior performance of a particular baseline algorithm over others.

```
[29]: # Create a dataframe to store the accuracy of baseline models for further
      ↪analysis
      basemodel_df = pd.DataFrame(
          columns=[
              "Baseline Model",
              "Accuracy",
              "Sensitivity",
              "Specificity",
              "False Positive Rate",
              "Precision",
              "F1-Score",
              "AUC Score"
          ]
      )
```

Logistic regression

Cross-Validation Set

```
[30]: with mlflow.start_run() as run:

      # Create a logistic regression classifier
      lr_clf_cv = LogisticRegression(random_state=42)

      # Fit the model to the training data
      lr_clf_cv.fit(X_train, y_train)

      # Make predictions on the test data
      y_pred_lr_cv = lr_clf_cv.predict(X_cv)
      # Get probabilities for the positive class
      y_scores_lr_cv = lr_clf_cv.predict_proba(X_cv)[: , 1]

      # Accuracy
      accuracy = lr_clf_cv.score(X_cv, y_cv)

      # Confusion Matrix
      cm = confusion_matrix(y_cv, y_pred_lr_cv)
      # Sensitivity (Recall) and Specificity calculations
      TP = cm[1, 1]
      TN = cm[0, 0]
      FP = cm[0, 1]
```

```

FN = cm[1, 0]
sensitivity = TP / (TP + FN)
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

# ROC AUC Score and curve
roc_auc = roc_auc_score(y_cv, y_scores_lr_cv) # Use scores, not
↪ predictions, for AUC
fpr, tpr, thresholds = roc_curve(y_cv, y_scores_lr_cv)

# Precision
precision = precision_score(y_cv, y_pred_lr_cv)

# F1-Score
f1 = f1_score(y_cv, y_pred_lr_cv)

# Append the new results to your DataFrame or storage structure
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "Logistic Regression CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Logistic Regression CV")
print(classification_report(y_cv, y_pred_lr_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(lr_clf_cv, "Logistic Regression CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(lr_clf_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,

```

```

        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Logistic Regression CV", ignore_errors=True)
#mlflow.sklearn.save_model(lr_clf_cv, "Logistic Regression CV")

signature = infer_signature(X_cv, y_pred_lr_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=lr_clf_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-logistic-reg-cv-model",
)

```

Logistic Regression CV

	precision	recall	f1-score	support
0	0.76	0.76	0.76	25
1	0.79	0.79	0.79	28
accuracy			0.77	53
macro avg	0.77	0.77	0.77	53
weighted avg	0.77	0.77	0.77	53

Confusion Matrix:

```
[[19  6]
 [ 6 22]]
```

Run ID: 82c8c30c68104370b012ef0b90298fc5

Registered model 'sk-learn-logistic-reg-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:32:23 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-logistic-reg-cv-model, version 2

Created version '2' of model 'sk-learn-logistic-reg-cv-model'.

```
[31]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_lr_cv)
```



```

roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_lr_cv)
pr_auc = average_precision_score(y_cv, y_scores_lr_cv)
f1 = f1_score(y_cv, y_pred_lr_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)

```

```

)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

# Update layout
fig.update_layout(
    title_text="Model Performance (Logistic Regression on Cross-Validation Set):  

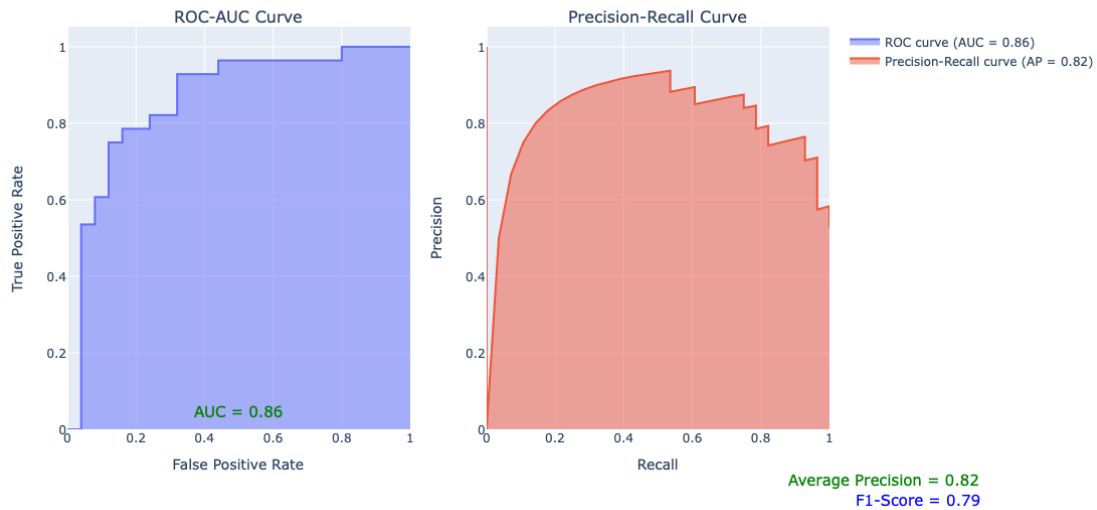
    ↳ ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (Logistic_
    ↳Regression on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
    ↳png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (Logistic_
    ↳Regression on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
    ↳svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (Logistic Regression on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

```
[32]: with mlflow.start_run() as run:

    # Create a logistic regression classifier
    lr_clf = LogisticRegression(random_state=42)

    # Fit the model to the training data
    lr_clf.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred_lr = lr_clf.predict(X_test)
    # Get probabilities for the positive class
    y_scores_lr = lr_clf.predict_proba(X_test)[: , 1]

    # Accuracy
    accuracy = lr_clf.score(X_test, y_test)

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred_lr)
    # Sensitivity (Recall) and Specificity calculations
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN)
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
```

```

FPR = FP / (FP + TN)

# ROC AUC Score and curve
roc_auc = roc_auc_score(y_test, y_scores_lr) # Use scores, not
↳ predictions, for AUC
fpr, tpr, thresholds = roc_curve(y_test, y_scores_lr)

# Precision
precision = precision_score(y_test, y_pred_lr)

# F1-Score
f1 = f1_score(y_test, y_pred_lr)

# Append the new results to your DataFrame or storage structure
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "Logistic Regression",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Logistic Regression")
print(classification_report(y_test, y_pred_lr))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(lr_clf, "Logistic Regression")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(lr_clf.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
    }
)

```

```

        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Logistic Regression", ignore_errors=True)
#mlflow.sklearn.save_model(lr_clf, "Logistic Regression")

signature = infer_signature(X_test, y_pred_lr)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=lr_clf,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-logistic-reg-model",
)

```

Logistic Regression

	precision	recall	f1-score	support
0	0.92	0.82	0.87	28
1	0.81	0.92	0.86	24
accuracy			0.87	52
macro avg	0.87	0.87	0.87	52
weighted avg	0.87	0.87	0.87	52

Confusion Matrix:

```
[[23  5]
 [ 2 22]]
```

Run ID: cabeda914b674295ac3ca39d71b49581

Registered model 'sk-learn-logistic-reg-model' already exists. Creating a new version of this model...

2024/04/13 12:32:32 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-logistic-reg-model, version 2

Created version '2' of model 'sk-learn-logistic-reg-model'.

```

[33]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_lr)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_lr)

```

```

pr_auc = average_precision_score(y_test, y_scores_lr)
f1 = f1_score(y_test, y_pred_lr)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,

```

```

        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)


# Update layout
fig.update_layout(
    title_text="Model Performance (Logistic Regression): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (Logistic Regression Set): ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (Logistic Regression Set): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (Logistic Regression): ROC-AUC and Precision-Recall Curves



SVM Classifier

Cross-Validation Set

```
[34]: with mlflow.start_run() as run:

    # Training the SVM model
    svm_model_cv = SVC(kernel="linear", probability=True)

    # Fit the model to the training data
    svm_model_cv.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred_svm_cv = svm_model_cv.predict(X_cv)
    # Get probabilities for the positive class
    y_scores_svm_cv = svm_model_cv.predict_proba(X_cv)[:, 1]

    # Accuracy
    accuracy = svm_model_cv.score(X_cv, y_cv)

    # Confusion Matrix
    cm = confusion_matrix(y_cv, y_pred_svm_cv)
    # Sensitivity (Recall) and Specificity calculations
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN)
```



```

specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

# ROC AUC Score and curve
roc_auc = roc_auc_score(y_cv, y_scores_svm_cv) # Use scores, not_
↪ predictions, for AUC
fpr, tpr, thresholds = roc_curve(y_cv, y_scores_svm_cv)

precision = precision_score(y_cv, y_pred_svm_cv)

# F1-Score
f1 = f1_score(y_cv, y_pred_svm_cv)

# Append the new results to your DataFrame or storage structure
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "SVM Classifier CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("SVM Classifier CV")
print(classification_report(y_cv, y_pred_svm_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(svm_model_cv, "SVM Classifier CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(svm_model_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
    }
)

```

```

        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("SVM Classifier CV", ignore_errors=True)
#mlflow.sklearn.save_model(svm_model_cv, "SVM Classifier CV")

signature = infer_signature(X_cv, y_pred_svm_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=svm_model_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-svm-clf-cv-model",
)

```

SVM Classifier CV

	precision	recall	f1-score	support
0	0.78	0.72	0.75	25
1	0.77	0.82	0.79	28
accuracy			0.77	53
macro avg	0.77	0.77	0.77	53
weighted avg	0.77	0.77	0.77	53

Confusion Matrix:

```
[[18  7]
 [ 5 23]]
```

Run ID: cc389d463ec5472297e23b7cf27c4c5d

Registered model 'sk-learn-svm-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:32:52 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-svm-clf-cv-model, version 2

Created version '2' of model 'sk-learn-svm-clf-cv-model'.

```

[35]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_svm_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve

```

```

precision, recall, _ = precision_recall_curve(y_cv, y_scores_svm_cv)
pr_auc = average_precision_score(y_cv, y_scores_svm_cv)
f1 = f1_score(y_cv, y_pred_svm_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,

```

```

        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

# Update layout
fig.update_layout(
    title_text="Model Performance (SVM Classifier on Cross-Validation Set):  

    ↳ ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (SVM Classifier on  

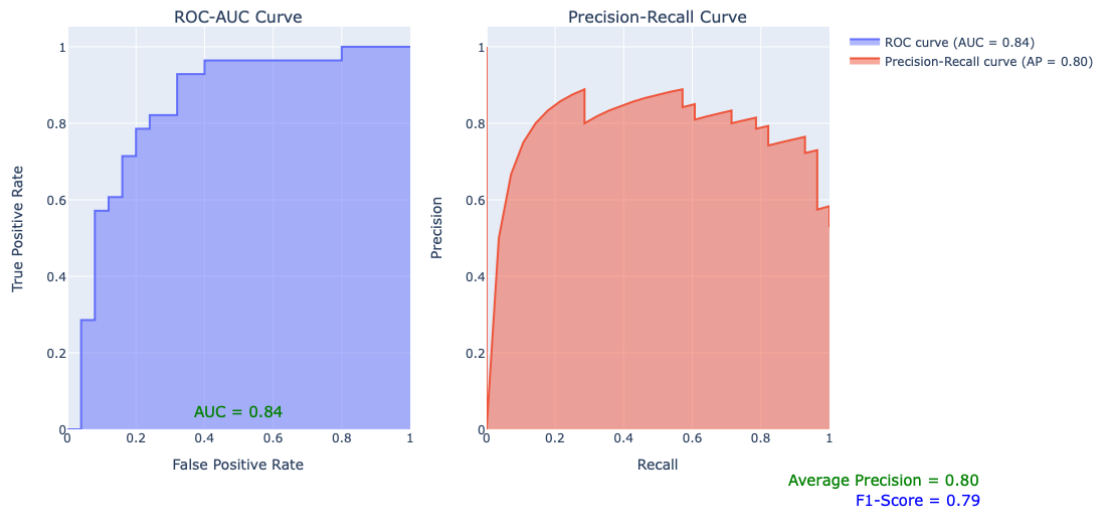
# ↳ Cross-Validation Set): ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (SVM Classifier on  

# ↳ Cross-Validation Set): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (SVM Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

```
[36]: with mlflow.start_run() as run:

    # Training the SVM model
    svm_model = SVC(kernel="linear", probability=True)

    # Fit the model to the training data
    svm_model.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred_svm = svm_model.predict(X_test)
    # Get probabilities for the positive class
    y_scores_svm = svm_model.predict_proba(X_test)[: , 1]

    # Accuracy
    accuracy = svm_model.score(X_test, y_test)

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred_svm)
    # Sensitivity (Recall) and Specificity calculations
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN)
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
```

```

FPR = FP / (FP + TN)

# ROC AUC Score and curve
roc_auc = roc_auc_score(y_test, y_scores_svm) # Use scores, not
↳ predictions, for AUC
fpr, tpr, thresholds = roc_curve(y_test, y_scores_svm)

precision = precision_score(y_test, y_pred_svm)

# F1-Score
f1 = f1_score(y_test, y_pred_svm)

# Append the new results to your DataFrame or storage structure
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "SVM Classifier",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("SVM Classifier")
print(classification_report(y_test, y_pred_svm))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(svm_model, "SVM Classifier")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(svm_model.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
    }
)

```

```

        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("SVM Classifier", ignore_errors=True)
#mlflow.sklearn.save_model(svm_model, "SVM Classifier")

signature = infer_signature(X_test, y_pred_svm)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=svm_model,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-svm-clf-model",
)

```

SVM Classifier

	precision	recall	f1-score	support
0	0.91	0.75	0.82	28
1	0.76	0.92	0.83	24
accuracy			0.83	52
macro avg	0.84	0.83	0.83	52
weighted avg	0.84	0.83	0.83	52

Confusion Matrix:

```

[[21  7]
 [ 2 22]]

```

Run ID: 3db929628a014493abf765583fa37704

Registered model 'sk-learn-svm-clf-model' already exists. Creating a new version of this model...

2024/04/13 12:33:13 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-svm-clf-model, version 2

Created version '2' of model 'sk-learn-svm-clf-model'.

```

[37]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_svm)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_svm)
pr_auc = average_precision_score(y_test, y_scores_svm)

```

```

f1 = f1_score(y_test, y_pred_svm)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",

```



```

    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

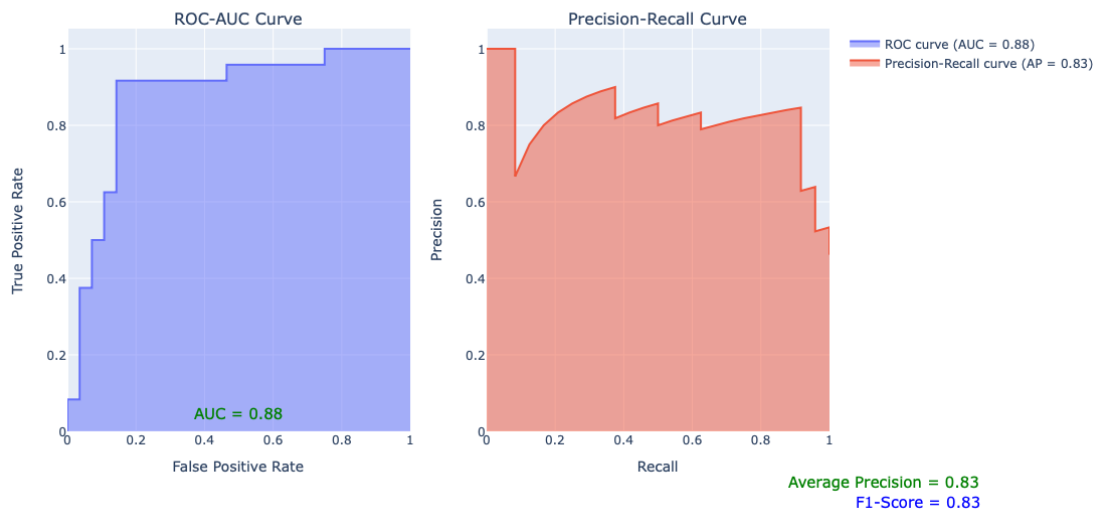
# Update layout
fig.update_layout(
    title_text="Model Performance (SVM Classifier): ROC-AUC and_
    ↳Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (SVM Classifier):_
    ↳ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (SVM Classifier):_
    ↳ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (SVM Classifier): ROC-AUC and Precision-Recall Curves



KNN Classifier

Cross-Validation Set

```
[38]: with mlflow.start_run() as run:

    # Training the KNN model
    knn_model_cv = KNeighborsClassifier(n_neighbors=5)

    # Fit the model to the training data
    knn_model_cv.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred_knn_cv = knn_model_cv.predict(X_cv)
    # Get probabilities for the positive class
    y_scores_knn_cv = knn_model_cv.predict_proba(X_cv)[:, 1]

    # Metrics calculation
    accuracy = knn_model_cv.score(X_cv, y_cv)
    cm = confusion_matrix(y_cv, y_pred_knn_cv)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)
```

```

roc_auc = roc_auc_score(y_cv, y_scores_knn_cv)
fpr, tpr, _ = roc_curve(y_cv, y_scores_knn_cv)
precision = precision_score(y_cv, y_pred_knn_cv)
f1 = f1_score(y_cv, y_pred_knn_cv)
pr_auc = average_precision_score(y_cv, y_scores_knn_cv)

# Append the new results to your DataFrame
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "KNN Classifier CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("KNN Classifier CV")
print(classification_report(y_cv, y_pred_knn_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(knn_model_cv, "KNN Classifier CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(knn_model_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow

```

```

#shutil.rmtree("KNN Classifier CV", ignore_errors=True)
#mlflow.sklearn.save_model(knn_model_cv, "KNN Classifier CV")

signature = infer_signature(X_cv, y_pred_knn_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=knn_model_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-knn-clf-cv-model",
)

```

KNN Classifier CV

	precision	recall	f1-score	support
0	0.70	0.76	0.73	25
1	0.77	0.71	0.74	28
accuracy			0.74	53
macro avg	0.74	0.74	0.74	53
weighted avg	0.74	0.74	0.74	53

Confusion Matrix:

```

[[19  6]
 [ 8 20]]

```

Run ID: 4017f0679d36477298e8e3b670d70203

Registered model 'sk-learn-knn-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:33:23 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-knn-clf-cv-model, version 3

Created version '3' of model 'sk-learn-knn-clf-cv-model'.

```

[39]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_knn_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_knn_cv)
pr_auc = average_precision_score(y_cv, y_scores_knn_cv)
f1 = f1_score(y_cv, y_pred_knn_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

```

```

)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)

```

```

fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

# Update layout
fig.update_layout(
    title_text="Model Performance (KNN Classifier on Cross-Validation Set):  

    ↳ ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (KNN Classifier on  


# ↳ Cross-Validation Set): ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (KNN Classifier on  

# ↳ Cross-Validation Set): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (KNN Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

```
[40]: with mlflow.start_run() as run:

    # Training the KNN model
    knn_model = KNeighborsClassifier(n_neighbors=5)

    # Fit the model to the training data
    knn_model.fit(X_train, y_train)

    # Make predictions on the test data
    y_pred_knn = knn_model.predict(X_test)
    # Get probabilities for the positive class
    y_scores_knn = knn_model.predict_proba(X_test)[:, 1]

    # Metrics calculation
    accuracy = knn_model.score(X_test, y_test)
    cm = confusion_matrix(y_test, y_pred_knn)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)
    roc_auc = roc_auc_score(y_test, y_scores_knn)
    fpr, tpr, _ = roc_curve(y_test, y_scores_knn)
```

```

precision = precision_score(y_test, y_pred_knn)
f1 = f1_score(y_test, y_pred_knn)
pr_auc = average_precision_score(y_test, y_scores_knn)

# Append the new results to your DataFrame
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "KNN Classifier",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("KNN Classifier")
print(classification_report(y_test, y_pred_knn))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(knn_model, "KNN Classifier")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(knn_model.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("KNN Classifier", ignore_errors=True)
#mlflow.sklearn.save_model(knn_model, "KNN Classifier")

```



```
signature = infer_signature(X_cv, y_pred_knn)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=knn_model,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-knn-clf-cv-model",
)
```

KNN Classifier

	precision	recall	f1-score	support
0	0.90	0.64	0.75	28
1	0.69	0.92	0.79	24
accuracy			0.77	52
macro avg	0.79	0.78	0.77	52
weighted avg	0.80	0.77	0.77	52

Confusion Matrix:

```
[[18 10]
 [ 2 22]]
```

Run ID: 25d0606e6c3e4df7ae58ad1f3401fd5b

Registered model 'sk-learn-knn-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:33:33 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-knn-clf-cv-model, version 4

Created version '4' of model 'sk-learn-knn-clf-cv-model'.

```
[41]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_knn)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_knn)
pr_auc = average_precision_score(y_test, y_scores_knn)
f1 = f1_score(y_test, y_pred_knn)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
```

```

fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,

```

```

    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)


# Update layout
fig.update_layout(
    title_text="Model Performance (KNN Classifier): ROC-AUC and_
↳Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (KNN Classifier):_
↳ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (KNN Classifier):_
↳ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (KNN Classifier): ROC-AUC and Precision-Recall Curves



Decision Tree Classifier

Cross-Validation Set

```
[42]: with mlflow.start_run() as run:

    # Create a random decision tree classifier
    dt_clf_cv = DecisionTreeClassifier(random_state=42)

    # Fit the model to the training data
    dt_clf_cv.fit(X_train, y_train)

    # Predict the test data
    y_pred_dt_cv = dt_clf_cv.predict(X_cv)
    # Get probabilities for the positive class
    y_scores_dt_cv = dt_clf_cv.predict_proba(X_cv)[:, 1]

    # Metrics calculation
    accuracy = accuracy_score(y_cv, y_pred_dt_cv)
    cm = confusion_matrix(y_cv, y_pred_dt_cv)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN)
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)
```

```

roc_auc = roc_auc_score(y_cv, y_scores_dt_cv)
fpr, tpr, _ = roc_curve(y_cv, y_scores_dt_cv)
precision = precision_score(y_cv, y_pred_dt_cv)
f1 = f1_score(y_cv, y_pred_dt_cv)
pr_auc = average_precision_score(y_cv, y_scores_dt_cv)

# Append the new results to your DataFrame
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "Decision Tree Classifier CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Decision Tree Classifier CV")
print(classification_report(y_cv, y_pred_dt_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(dt_clf_cv, "Decision Tree Classifier CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(dt_clf_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow

```

```

#shutil.rmtree("Decision Tree Classifier CV", ignore_errors=True)
#mlflow.sklearn.save_model(dt_clf_cv, "Decision Tree Classifier CV")

signature = infer_signature(X_cv, y_pred_dt_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=dt_clf_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-decision-tree-clf-cv-model",
)

```

Decision Tree Classifier CV

	precision	recall	f1-score	support
0	0.79	0.76	0.78	25
1	0.79	0.82	0.81	28
accuracy			0.79	53
macro avg	0.79	0.79	0.79	53
weighted avg	0.79	0.79	0.79	53

Confusion Matrix:

```
[[19  6]
 [ 5 23]]
```

Run ID: b80733af1ea341f3b574f35f1c81bb08

Registered model 'sk-learn-decision-tree-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:33:42 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-decision-tree-clf-cv-model, version 3

Created version '3' of model 'sk-learn-decision-tree-clf-cv-model'.

```

[43]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_dt_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_dt_cv)
pr_auc = average_precision_score(y_cv, y_scores_dt_cv)
f1 = f1_score(y_cv, y_pred_dt_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

```

```

)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)

```

```

fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)


# Update layout
fig.update_layout(
    title_text="Model Performance (Decision Tree Classifier on Cross-Validation_
↪Set): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Baseline_Models/Model Performance (Decision Tree_
↪Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
↪png") #png format
#fig.write_image("./charts/Baseline_Models/Model Performance (Decision Tree_
↪Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
↪svg") #svg format

# Display the plots side-by-side
fig.show()

```


Model Performance (Decision Tree Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

```
[44]: with mlflow.start_run() as run:

    # Create a random decision tree classifier
    dt_clf = DecisionTreeClassifier(random_state=42)

    # Fit the model to the training data
    dt_clf.fit(X_train, y_train)

    # Predict on the test data
    y_pred_dt = dt_clf.predict(X_test)
    # Get probabilities for the positive class
    y_scores_dt = dt_clf.predict_proba(X_test)[:, 1]

    # Metrics calculation
    accuracy = accuracy_score(y_test, y_pred_dt)
    cm = confusion_matrix(y_test, y_pred_dt)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN)
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)

    roc_auc = roc_auc_score(y_test, y_scores_dt)
```

```

fpr, tpr, _ = roc_curve(y_test, y_scores_dt)
precision = precision_score(y_test, y_pred_dt)
f1 = f1_score(y_test, y_pred_dt)
pr_auc = average_precision_score(y_test, y_scores_dt)

# Append the new results to your DataFrame
basemodel_df = basemodel_df._append(
    {
        "Baseline Model": "Decision Tree Classifier",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Decision Tree Classifier")
print(classification_report(y_test, y_pred_dt))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(dt_clf, "Decision Tree Classifier")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(dt_clf.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Decision Tree Classifier", ignore_errors=True)
#mlflow.sklearn.save_model(dt_clf, "Decision Tree Classifier")

```

```

signature = infer_signature(X_test, y_pred_dt_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=dt_clf,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-decision-tree-clf-cv-model",
)

```

Decision Tree Classifier

	precision	recall	f1-score	support
0	0.71	0.71	0.71	28
1	0.67	0.67	0.67	24
accuracy			0.69	52
macro avg	0.69	0.69	0.69	52
weighted avg	0.69	0.69	0.69	52

Confusion Matrix:

```

[[20  8]
 [ 8 16]]

```

Run ID: b6966e470e9e44d08ebcb04e11f1ff46

Registered model 'sk-learn-decision-tree-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:33:51 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-decision-tree-clf-cv-model, version 4

Created version '4' of model 'sk-learn-decision-tree-clf-cv-model'.

```

[45]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_dt)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_dt)
pr_auc = average_precision_score(y_test, y_scores_dt)
f1 = f1_score(y_test, y_pred_dt)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

```

```

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,

```

```

        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Decision Tree Classifier): ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(
        margin=dict(b=100)
    ) # Adjust bottom margin to avoid cutting off annotations

    # Save plot
    #fig.write_image("./charts/Baseline_Models/Model Performance (Decision Tree Classifier): ROC-AUC and Precision-Recall Curves.png") #png format
    #fig.write_image("./charts/Baseline_Models/Model Performance (Decision Tree Classifier): ROC-AUC and Precision-Recall Curves.svg") #svg format

    # Display the plots side-by-side
    fig.show()

```

Model Performance (Decision Tree Classifier): ROC-AUC and Precision-Recall Curves



Evaluation Metrics Comparison

```
[46]: basemodel_df.sort_values("Specificity", ascending=False)
```

```
[46]:
```

	Baseline Model	Accuracy	Sensitivity	Specificity	False
Positive Rate	Precision	F1-Score	AUC Score		
1	Logistic Regression	0.865385	0.916667	0.821429	
0.178571	0.814815	0.862745	0.882440		
0	Logistic Regression CV	0.773585	0.785714	0.760000	
0.240000	0.785714	0.785714	0.862857		
6	Decision Tree Classifier CV	0.792453	0.821429	0.760000	
0.240000	0.793103	0.807018	0.790714		
4	KNN Classifier CV	0.735849	0.714286	0.760000	
0.240000	0.769231	0.740741	0.785000		
3	SVM Classifier	0.826923	0.916667	0.750000	
0.250000	0.758621	0.830189	0.875000		
2	SVM Classifier CV	0.773585	0.821429	0.720000	
0.280000	0.766667	0.793103	0.844286		
7	Decision Tree Classifier	0.692308	0.666667	0.714286	
0.285714	0.666667	0.666667	0.690476		
5	KNN Classifier	0.769231	0.916667	0.642857	
0.357143	0.687500	0.785714	0.837054		

The *Logistic Regression algorithm* provided the best performance among the four (4) baseline models in terms of all evaluation metrics used.

Performance Metrics of Baseline Models

```

[47]: # Define a list of colors for the baseline models
colors = [
    "rgba(255, 99, 132, 0.6)",
    "rgba(54, 162, 235, 0.6)",
    "rgba(255, 206, 86, 0.6)",
    "rgba(75, 192, 192, 0.6)",
    "rgba(153, 102, 255, 0.6)",
    "rgba(255, 159, 64, 0.6)",
]

# Map each ensemble model to a specific color
unique_models = basemodel_df["Baseline Model"].unique()
color_map = {model: colors[i % len(colors)] for i, model in enumerate(unique_models)}

# Create Subplots for the model performance using the 6 evaluation metrics
fig = make_subplots(
    rows=4,
    cols=2,
    subplot_titles=(
        "Accuracy vs. Baseline Model",
        "Specificity vs. Baseline Model",
        "Sensitivity vs. Baseline Model",
        "False Positive Rate vs. Baseline Model",
        "Precision vs. Baseline Model",
        "F1-Score vs. Baseline Model",
        "AUC Score vs. Baseline Model",
    ),
    horizontal_spacing=0.15,
    vertical_spacing=0.15,
)

metrics = [
    "Accuracy",
    "Specificity",
    "Sensitivity",
    "False Positive Rate",
    "Precision",
    "F1-Score",
    "AUC Score",
]

plot_positions = [(1, 1), (1, 2), (2, 1), (2, 2), (3, 1), (3, 2), (4, 1)]

for metric, pos in zip(metrics, plot_positions):
    # Sort the DataFrame based on the current metric in descending order
    df_sorted = basemodel_df.sort_values(by=metric, ascending=False)

```

```

# Extracting the sorted model names for consistent color mapping
sorted_models = df_sorted["Baseline Model"].unique()

# Generate one bar for each model, now in sorted order
for model in sorted_models:
    df_filtered = df_sorted[df_sorted["Baseline Model"] == model]
    show_legend = (
        metric == "Accuracy"
    ) # Show legend only in the first subplot for clarity
    fig.add_trace(
        go.Bar(
            x=[model],
            y=df_filtered[metric],
            name=model,
            marker_color=color_map[model],
            text=df_filtered[metric].round(2),
            textposition="outside",
            showlegend=show_legend,
        ),
        row=pos[0],
        col=pos[1],
    )

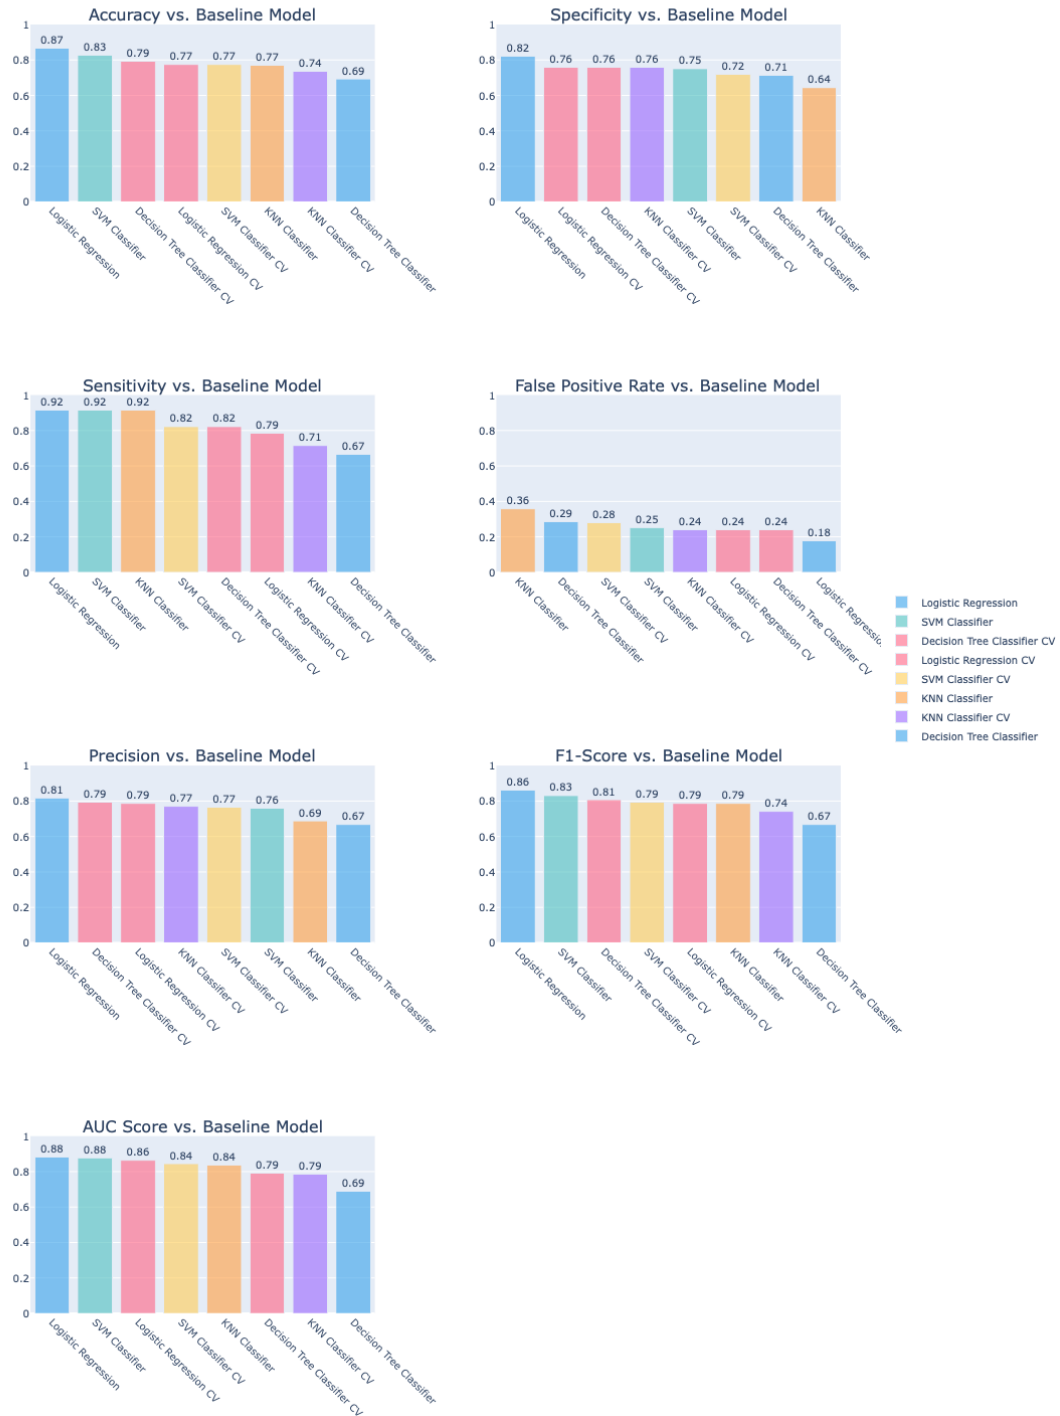
# Update layout
fig.update_layout(
    height=1500,
    width=1100,
    title_text="Performance Metrics of Baseline Models",
    showlegend=True,
    legend=dict(orientation="v", x=1.05, y=0.5),
    font=dict(size=10),
)
fig.update_xaxes(tickangle=45)
fig.update_yaxes(range=[0, 1])

# Save plot
#fig.write_image("./charts/Baseline_Models/Performance Metrics of Baseline_
↳Models.png") #png format
#fig.write_image("./charts/Baseline_Models/Performance Metrics of Baseline_
↳Models.svg") #svg format

# Display plot
fig.show()

```


Performance Metrics of Baseline Models



The **Logistic Regression algorithm** provided the best performance among the four baseline models in terms of all evaluation metrics used. The evaluation metrics used for comparison were **Accuracy**, **Sensitivity**, **Specificity**, **Precision**, **F1-Score**, and **AUC Score** respectively.

1.1.7 Analysis of Ensemble Methods

This section of the code involves evaluating the performance of the ensemble learning methods used in the model training section, by using appropriate evaluation metrics and comparing the results. The purpose is to determine the most effective and efficient method for detecting ovarian cancer tumors and to identify the factors that contribute to the superior performance of a particular ensemble learning technique over others.

```
[48]: # Create a dataframe to store the accuracy of ensemble models for further
      ↪analysis
ensemble_df = pd.DataFrame(
    columns=[
        "Ensemble Model",
        "Accuracy",
        "Sensitivity",
        "Specificity",
        "False Positive Rate",
        "Precision",
        "AUC Score",
        "F1-Score",
    ]
)
```

Voting Classifier

Cross-Validation Set

```
[49]: with mlflow.start_run() as run:

    # Making the final model using a voting classifier with soft voting
    vote_model_soft_cv = VotingClassifier(
        estimators=[
            ("logistic regression CV", lr_clf_cv),
            ("svc CV", svm_model_cv), # Make sure svm_model is trained with
            ↪probabilities
            ("knn CV", knn_model_cv),
            ("decision tree CV", dt_clf_cv),
        ],
        voting="soft",
    )

    # Training the model on the training dataset
    vote_model_soft_cv.fit(X_train, y_train)

    # Predicting the output on the test dataset
    y_pred_final_vm_soft_cv = vote_model_soft_cv.predict(X_cv)
    # Get probabilities for the positive class for AUC calculation
    y_scores_vm_soft_cv = vote_model_soft_cv.predict_proba(X_cv)[: , 1]
```

```

# Metrics calculation
accuracy = accuracy_score(y_cv, y_pred_final_vm_soft_cv)
cm = confusion_matrix(y_cv, y_pred_final_vm_soft_cv)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN)
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)
precision = precision_score(y_cv, y_pred_final_vm_soft_cv)
f1 = f1_score(y_cv, y_pred_final_vm_soft_cv)
roc_auc = roc_auc_score(y_cv, y_scores_vm_soft_cv)
fpr, tpr, _ = roc_curve(y_cv, y_scores_vm_soft_cv)

ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Voting Classifier (Soft) CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Voting Classifier Model (Soft Voting) CV")
print(classification_report(y_cv, y_pred_final_vm_soft_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(vote_model_soft_cv, "Voting Classifier Model (Soft_
↳Voting) Regression CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(vote_model_soft_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,

```

```

        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Voting Classifier Model (Soft Voting) CV",
↳ ignore_errors=True)
#mlflow.sklearn.save_model(vote_model_soft_cv, "Voting Classifier Model
↳ (Soft Voting) CV")

signature = infer_signature(X_cv, y_pred_final_vm_soft_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=vote_model_soft_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-svm-clf-cv-model",
)

```

Voting Classifier Model (Soft Voting) CV					
	precision	recall	f1-score	support	
0	0.79	0.76	0.78	25	
1	0.79	0.82	0.81	28	
accuracy			0.79	53	
macro avg	0.79	0.79	0.79	53	
weighted avg	0.79	0.79	0.79	53	

Confusion Matrix:

```
[[19  6]
 [ 5 23]]
```

Run ID: 94a7592ef594451ca382d59730822b91

Registered model 'sk-learn-svm-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:34:10 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-svm-clf-cv-model, version 5

Created version '5' of model 'sk-learn-svm-clf-cv-model'.

```
[50]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_vm_soft_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_vm_soft_cv)
pr_auc = average_precision_score(y_cv, y_scores_vm_soft_cv)
f1 = f1_score(y_cv, y_pred_final_vm_soft_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
```

```

        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Voting Classifier on Cross-Validation Set):  

        ↳ ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(
        margin=dict(b=100)
    )
    # Adjust bottom margin to avoid cutting off annotations

    # Save plot
    #fig.write_image("./charts/Ensemble_Models/Model Performance (Voting Classifier  

    ↳ on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.png") #png  

    ↳ format
    #fig.write_image("./charts/Ensemble_Models/Model Performance (Voting Classifier  

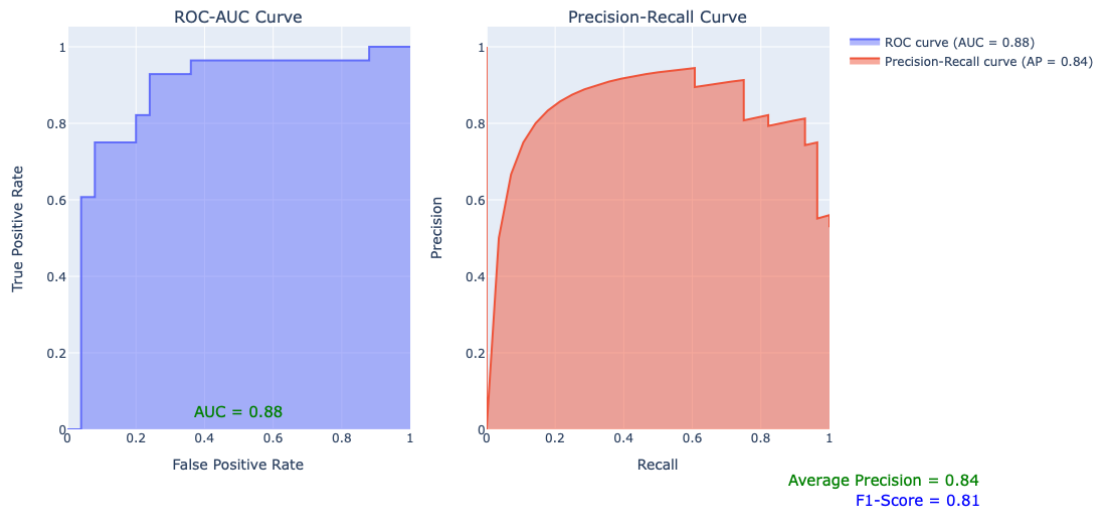
    ↳ on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.svg") #svg  

    ↳ format

    # Display the plots side-by-side
    fig.show()

```

Model Performance (Voting Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

[51]: `with mlflow.start_run() as run:`

```
# Making the final model using a voting classifier with soft voting
vote_model_soft = VotingClassifier(
    estimators=[
        ("logistic regression", lr_clf),
        ("svc", svm_model), # Make sure svm_model is trained with
        probabilities
        ("knn", knn_model),
        ("decision tree", dt_clf),
    ],
    voting="soft",
)

# Training the model on the training dataset
vote_model_soft.fit(X_train, y_train)

# Predicting the output on the test dataset
y_pred_final_vm_soft = vote_model_soft.predict(X_test)
# Get probabilities for the positive class for AUC calculation
y_scores_vm_soft = vote_model_soft.predict_proba(X_test)[:, 1]

# Metrics calculation
accuracy = accuracy_score(y_test, y_pred_final_vm_soft)
cm = confusion_matrix(y_test, y_pred_final_vm_soft)
TP = cm[1, 1]
```

```

TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN)
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

precision = precision_score(y_test, y_pred_final_vm_soft)
f1 = f1_score(y_test, y_pred_final_vm_soft)
roc_auc = roc_auc_score(y_test, y_scores_vm_soft)
fpr, tpr, _ = roc_curve(y_test, y_scores_vm_soft)

ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Voting Classifier (Soft)",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Voting Classifier Model (Soft Voting)")
print(classification_report(y_test, y_pred_final_vm_soft))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(vote_model_soft, "Voting Classifier Model (Soft_
↳Voting) CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(vote_model_soft.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
    }
)

```



```

        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Voting Classifier Model (Soft Voting) CV", ignore_errors=True)
mlflow.sklearn.save_model(vote_model_soft, "Voting Classifier Model (Soft Voting) CV")

signature = infer_signature(X_test, y_pred_final_vm_soft)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=vote_model_soft,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-voting-clf-cv-model",
)

```

Voting Classifier Model (Soft Voting)

	precision	recall	f1-score	support
0	0.96	0.86	0.91	28
1	0.85	0.96	0.90	24
accuracy			0.90	52
macro avg	0.91	0.91	0.90	52
weighted avg	0.91	0.90	0.90	52

Confusion Matrix:

```
[[24  4]
 [ 1 23]]
```

Run ID: ebaaad36eb7846ce8b31216d6a7c58ab

Registered model 'sk-learn-voting-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:34:28 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-voting-clf-cv-model, version 2

Created version '2' of model 'sk-learn-voting-clf-cv-model'.

[52]: *# Calculate metrics for ROC-AUC Curve*

```

fpr, tpr, _ = roc_curve(y_test, y_scores_vm_soft)
roc_auc_val = auc(fpr, tpr)

```

```

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_vm_soft)
pr_auc = average_precision_score(y_test, y_scores_vm_soft)
f1 = f1_score(y_test, y_pred_final_vm_soft)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tp,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(

```

```

        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

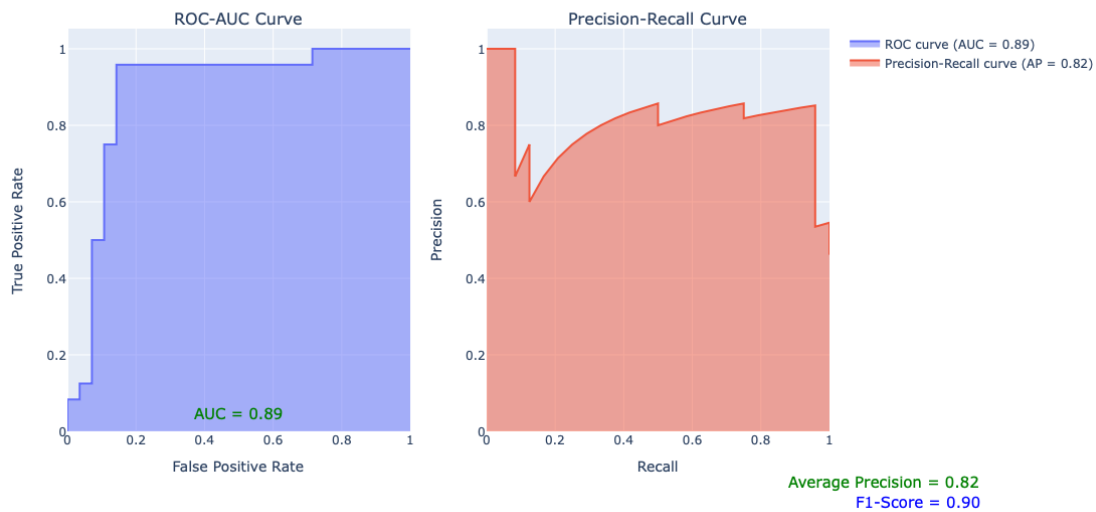
# Update layout
fig.update_layout(
    title_text="Model Performance (Voting Classifier): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust bottom margin to avoid cutting off annotations

# Save plot
fig.write_image("./charts/Ensemble_Models/Model Performance (Voting Classifier): ROC-AUC and Precision-Recall Curves.png") #png format
fig.write_image("./charts/Ensemble_Models/Model Performance (Voting Classifier): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (Voting Classifier): ROC-AUC and Precision-Recall Curves



Bagging Classifier

Cross-Validation Set

[53]: `with mlflow.start_run() as run:`

```
# Initializing the bagging model using SVM as the base model with default
parameters
bag_model_cv = BaggingClassifier(
    svm_model, random_state=42, n_estimators=10
)

# Training the model
bag_model_cv.fit(X_train, y_train)

# Predicting the output on the test dataset
y_pred_bm_cv = bag_model_cv.predict(X_cv)

# Calculate metrics
accuracy = accuracy_score(y_cv, y_pred_bm_cv)
cm = confusion_matrix(y_cv, y_pred_bm_cv)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
```

```

FPR = FP / (FP + TN)

precision = precision_score(y_cv, y_pred_bm_cv)
f1 = f1_score(y_cv, y_pred_bm_cv)

# Attempt to calculate the ROC AUC Score if probabilities can be estimated
try:
    y_scores_bm_cv = bag_model_cv.predict_proba(X_cv)[: , 1]
    roc_auc = roc_auc_score(y_cv, y_scores_bm_cv)
except AttributeError as e:
    roc_auc = 0.00
    print(
        "ROC AUC Score is not applicable for this configuration without_
↪predict_proba:",
        e,
    )

# Update ensemble_df DataFrame
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Bagging Classifier with SVM CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the validation set
print("Bagging Classifier with SVM CV: ")
print(classification_report(y_cv, y_pred_bm_cv))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(bag_model_cv, "Bagging Classifier with SVM CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(bag_model_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,

```

```

        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Bagging Classifier with SVM CV", ignore_errors=True)
#mlflow.sklearn.save_model(bag_model_cv, "Bagging Classifier with SVM CV")

signature = infer_signature(X_cv, y_pred_bm_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=bag_model_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-bagging-clf-cv-model",
)

```

Bagging Classifier with SVM CV:

	precision	recall	f1-score	support
0	0.88	0.60	0.71	25
1	0.72	0.93	0.81	28
accuracy			0.77	53
macro avg	0.80	0.76	0.76	53
weighted avg	0.80	0.77	0.77	53

Confusion Matrix:

```
[[15 10]
 [ 2 26]]
```

Run ID: 444d85fe6d3c4573b5f8f198bdf41a32

Registered model 'sk-learn-bagging-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:35:33 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-bagging-clf-cv-model, version 3

Created version '3' of model 'sk-learn-bagging-clf-cv-model'.

```
[54]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_bm_cv)
roc_auc_val = auc(fpr, tpr)
```

```

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_bm_cv)
pr_auc = average_precision_score(y_cv, y_scores_bm_cv)
f1 = f1_score(y_cv, y_pred_bm_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(

```

```

        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

# Update layout
fig.update_layout(
    title_text="Model Performance (Bagging Classifier on Cross-Validation Set):  

    ↳ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
fig.write_image("./charts/Ensemble_Models/Model Performance (Bagging  

    ↳Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.  

    ↳png") #png format
fig.write_image("./charts/Ensemble_Models/Model Performance (Bagging  


    ↳Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.  

    ↳svg") #svg format

# Display the plots side-by-side
fig.show()

```


Model Performance (Bagging Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

```
[55]: with mlflow.start_run() as run:
```

```
    # Initializing the bagging model using SVM as the base model with default_
    ↪ parameters
    bag_model = BaggingClassifier(
        svm_model, random_state=42, n_estimators=10
    )

    # Training the model
    bag_model.fit(X_train, y_train)

    # Predicting the output on the test dataset
    y_pred_bm = bag_model.predict(X_test)

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred_bm)
    cm = confusion_matrix(y_test, y_pred_bm)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)
```

```

precision = precision_score(y_test, y_pred_bm)
f1 = f1_score(y_test, y_pred_bm)

# Attempt to calculate the ROC AUC Score if probabilities can be estimated
try:
    y_scores_bm = bag_model.predict_proba(X_test)[: , 1]
    roc_auc = roc_auc_score(y_test, y_scores_bm)
except AttributeError as e:
    roc_auc = 0.00
    print("ROC AUC Score is not applicable for this configuration without_
↪predict_proba:", e)

# Update ensemble_df DataFrame
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Bagging Classifier with SVM",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Bagging Classifier with SVM: ")
print(classification_report(y_test, y_pred_bm))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(bag_model, "Bagging Classifier with SVM")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(bag_model.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

```

```

    }
)

# Save the model to MLflow
#shutil.rmtree("Bagging Classifier with SVM", ignore_errors=True)
#mlflow.sklearn.save_model(bag_model, "Bagging Classifier with SVM")

signature = infer_signature(X_cv, y_pred_bm)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=bag_model,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-logistic-reg-cv-model",
)

```

Bagging Classifier with SVM:

	precision	recall	f1-score	support
0	0.91	0.75	0.82	28
1	0.76	0.92	0.83	24
accuracy			0.83	52
macro avg	0.84	0.83	0.83	52
weighted avg	0.84	0.83	0.83	52

Confusion Matrix:

```
[[21  7]
 [ 2 22]]
```

Run ID: a22c126e92f84badad5de59e6620af5f

Registered model 'sk-learn-logistic-reg-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:36:39 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-logistic-reg-cv-model, version 6

Created version '6' of model 'sk-learn-logistic-reg-cv-model'.

```

[56]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_bm)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_bm)
pr_auc = average_precision_score(y_test, y_scores_bm)
f1 = f1_score(y_test, y_pred_bm)

```

```

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tp,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozero",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozero",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",

```

```

        showarrow=False,
        font=dict(size=15, color="green"),
    )
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)


# Update layout
fig.update_layout(
    title_text="Model Performance (Bagging Classifier): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
fig.write_image("./charts/Ensemble_Models/Model Performance (Bagging Classifier): ROC-AUC and Precision-Recall Curves.png") #png format
fig.write_image("./charts/Ensemble_Models/Model Performance (Bagging Classifier): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (Bagging Classifier): ROC-AUC and Precision-Recall Curves



GBM Classifier

Cross-Validation Set

[57]: `with mlflow.start_run() as run:`

```
# Initializing the Gradient Boosting classifier with default parameters
gb_model_cv = GradientBoostingClassifier()

# Training the model on the training dataset
gb_model_cv.fit(X_train, y_train)

# Predicting the output on the test dataset
y_pred_final_gb_cv = gb_model_cv.predict(X_cv)
# Get probability scores for AUC calculation
y_scores_gb_cv = gb_model_cv.predict_proba(X_cv)[: , 1]

# Calculate metrics
accuracy = accuracy_score(y_cv, y_pred_final_gb_cv)
cm = confusion_matrix(y_cv, y_pred_final_gb_cv)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)
```

```

precision = precision_score(y_cv, y_pred_final_gb_cv)
f1 = f1_score(y_cv, y_pred_final_gb_cv)
roc_auc = roc_auc_score(y_cv, y_scores_gb_cv) # Use probability scores for AUC
↪AUC

# Update ensemble_df DataFrame with new metrics
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Gradient Boosting Classifier CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the validation set
print("Gradient Boosting Classifier CV: ")
print(classification_report(y_cv, y_pred_final_gb_cv))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(gb_model_cv, "Gradient Boosting Classifier CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(gb_model_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Stacking Ensemble Models CV", ignore_errors=True)
#mlflow.sklearn.save_model(gb_model_cv, "Stacking Ensemble Models CV")

```

```
signature = infer_signature(X_cv, y_pred_final_gb_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=gb_model_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-gbm-clf-cv-model",
)
```

Gradient Boosting Classifier CV:

	precision	recall	f1-score	support
0	0.87	0.80	0.83	25
1	0.83	0.89	0.86	28
accuracy			0.85	53
macro avg	0.85	0.85	0.85	53
weighted avg	0.85	0.85	0.85	53

Confusion Matrix:

```
[[20  5]
 [ 3 25]]
```

Run ID: 7f773c00937f4c66a8f634bfdd580629

Registered model 'sk-learn-gbm-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:36:48 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-gbm-clf-cv-model, version 3

Created version '3' of model 'sk-learn-gbm-clf-cv-model'.

```
[58]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_gb_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_gb_cv)
pr_auc = average_precision_score(y_cv, y_scores_gb_cv)
f1 = f1_score(y_cv, y_pred_final_gb_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
```



```

fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,

```

```

xref="paper",
yref="paper",
text=f"F1-Score = {f1:.2f}",
showarrow=False,
font=dict(size=15, color="blue"),
)

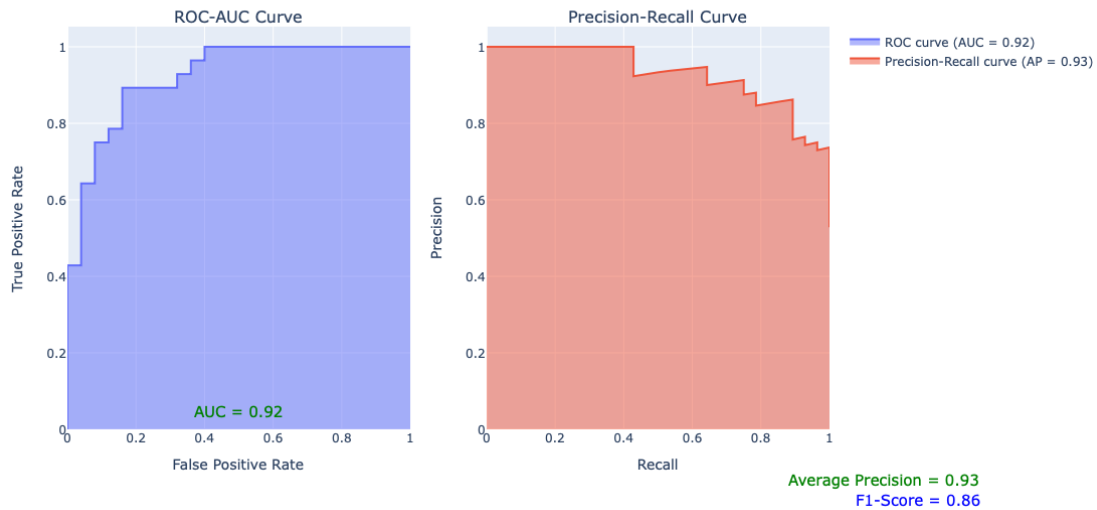
# Update layout
fig.update_layout(
    title_text="Model Performance (Gradient Boosting Classifier on_
↳Cross-Validation Set): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Gradient Boosting_
↳Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
↳png") #png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Gradient Boosting_
↳Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves.
↳svg") #svg format

# Display the plots side-by-side
fig.show()

```

Model Performance (Gradient Boosting Classifier on Cross-Validation Set): ROC-AUC and Precision-Recall Curves



Testing Set

[59]: `with mlflow.start_run() as run:`

```
# Initializing the Gradient Boosting classifier with default parameters
gb_model = GradientBoostingClassifier()

# Training the model on the training dataset
gb_model.fit(X_train, y_train)

# Predicting the output on the test dataset
y_pred_final_gb = gb_model.predict(X_test)
# Get probability scores for AUC calculation
y_scores_gb = gb_model.predict_proba(X_test)[: , 1]

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred_final_gb)
cm = confusion_matrix(y_test, y_pred_final_gb)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

precision = precision_score(y_test, y_pred_final_gb)
```

```

f1 = f1_score(y_test, y_pred_final_gb)
roc_auc = roc_auc_score(y_test, y_scores_gb) # Use probability scores for AUC
↪AUC

# Update ensemble_df DataFrame with new metrics
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Gradient Boosting Classifier",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("Gradient Boosting Classifier: ")
print(classification_report(y_test, y_pred_final_gb))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model.gb_model, "Gradient Boosting Classifier")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params.gb_model.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Gradient Boosting Classifier", ignore_errors=True)
#mlflow.sklearn.save_model.gb_model, "Gradient Boosting Classifier")

signature = infer_signature(X_cv, y_pred_final_gb)

```

```

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=gb_model,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-gbm-clf-model",
)

```

Gradient Boosting Classifier:

	precision	recall	f1-score	support
0	0.92	0.86	0.89	28
1	0.85	0.92	0.88	24
accuracy			0.88	52
macro avg	0.88	0.89	0.88	52
weighted avg	0.89	0.88	0.88	52

Confusion Matrix:

```

[[24  4]
 [ 2 22]]

```

Run ID: 88a96c9d1a1f4eca8e20ea50e7107d26

Registered model 'sk-learn-gbm-clf-model' already exists. Creating a new version of this model...

2024/04/13 12:36:56 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-gbm-clf-model, version 3

Created version '3' of model 'sk-learn-gbm-clf-model'.

```

[60]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_gb)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_gb)
pr_auc = average_precision_score(y_test, y_scores_gb)
f1 = f1_score(y_test, y_pred_final_gb)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(

```

```

        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)
fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",

```

```

text=f"F1-Score = {f1:.2f}",
showarrow=False,
font=dict(size=15, color="blue"),
)

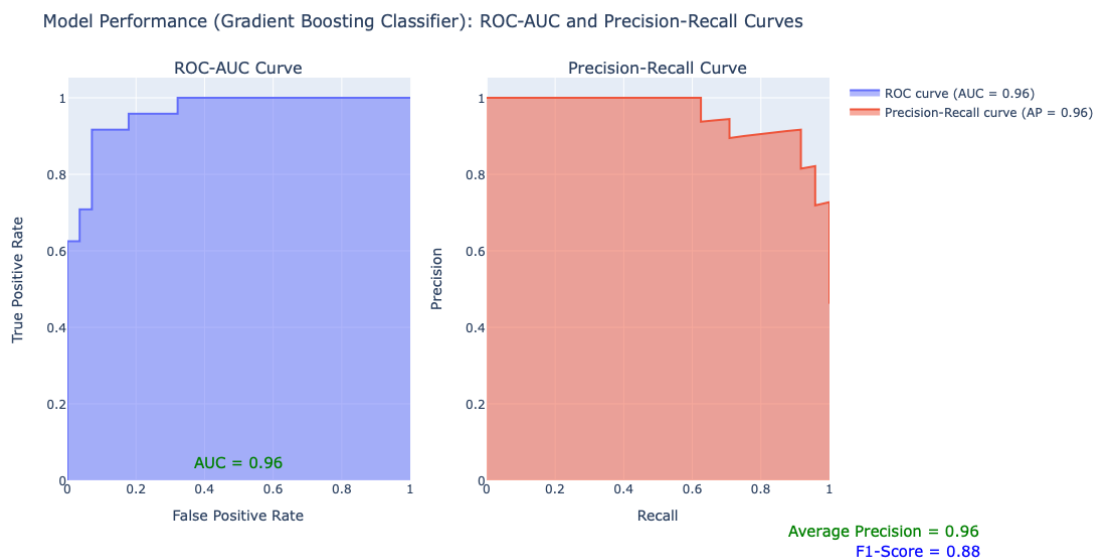
# Update layout
fig.update_layout(
    title_text="Model Performance (Gradient Boosting Classifier): ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)

fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Gradient Boosting Classifier): ROC-AUC and Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Gradient Boosting Classifier): ROC-AUC and Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```



XGBoost Classifier

Cross-Validation Set

```
[61]: with mlflow.start_run() as run:

    # Initializing the XGBoost classifier with default parameters
    xgb_model_cv = XGBClassifier(use_label_encoder=False, eval_metric="logloss")

    # Training the model on the train dataset
    xgb_model_cv.fit(X_train, y_train)

    # Predicting the output on the test dataset
    y_pred_final_xgb_cv = xgb_model_cv.predict(X_cv)
    # Get probability scores for AUC calculation
    y_scores_xgb_cv = xgb_model_cv.predict_proba(X_cv)[:, 1]

    # Calculate metrics
    accuracy = accuracy_score(y_cv, y_pred_final_xgb_cv)
    cm = confusion_matrix(y_cv, y_pred_final_xgb_cv)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)

    precision = precision_score(y_cv, y_pred_final_xgb_cv)
    f1 = f1_score(y_cv, y_pred_final_xgb_cv)
    roc_auc = roc_auc_score(y_cv, y_scores_xgb_cv) # Use probability scores
    ↪ for AUC

    # Update ensemble_df DataFrame
    ensemble_df = ensemble_df._append(
        {
            "Ensemble Model": "XGBoost Classifier CV",
            "Accuracy": accuracy,
            "Sensitivity": sensitivity,
            "Specificity": specificity,
            "False Positive Rate": FPR,
            "AUC Score": roc_auc,
            "F1-Score": f1,
            "Precision": precision,
        },
```



```

        ignore_index=True,
    )

    # Evaluate the model on the validation set
    print("XGBoost Classifier CV: ")
    print(classification_report(y_cv, y_pred_final_xgb_cv))
    print("Confusion Matrix: ")
    print(cm)

    # Log the model parameters and metrics to MLflow
    mlflow.sklearn.log_model(xgb_model_cv, "XGBoost Classifier CV")
    print("Run ID: {}".format(run.info.run_id))

    mlflow.log_params(xgb_model_cv.get_params())
    mlflow.log_metrics(
        {
            "Accuracy": accuracy,
            "Sensitivity": sensitivity,
            "Specificity": specificity,
            "False Positive Rate": FPR,
            "AUC Score": roc_auc,
            "F1-Score": f1,
            "Precision": precision
        }
    )

    # Save the model to MLflow
    #shutil.rmtree("XGBoost Classifier CV", ignore_errors=True)
    #mlflow.xgboost.save_model(xgb_model_cv, "XGBoost Classifier CV")

    signature = infer_signature(X_cv, y_pred_final_xgb_cv)

    # Log the sklearn model and register as version 1
    mlflow.xgboost.log_model(
        xgb_model=xgb_model_cv,
        artifact_path="xgboost-model",
        signature=signature,
        registered_model_name="xgboost-clf-cv-model",
    )

```

XGBoost Classifier CV:

	precision	recall	f1-score	support
0	0.83	0.80	0.82	25
1	0.83	0.86	0.84	28
accuracy			0.83	53
macro avg	0.83	0.83	0.83	53

weighted avg 0.83 0.83 0.83 53

Confusion Matrix:

```
[[20  5]
 [ 4 24]]
```

Run ID: 1eebe7d988e04352b4484c82d13c3301

Registered model 'xgboost-clf-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:37:05 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: xgboost-clf-cv-model, version 3

Created version '3' of model 'xgboost-clf-cv-model'.

```
[62]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_xgb_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_xgb_cv)
pr_auc = average_precision_score(y_cv, y_scores_xgb_cv)
f1 = f1_score(y_cv, y_pred_final_xgb_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
```

```

        row=1,
        col=1,
    )

    # Add Precision-Recall Curve to the subplot
    fig.add_trace(
        go.Scatter(
            x=recall,
            y=precision,
            mode="lines",
            name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
            fill="tozeroy",
        ),
        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Extreme Gradient Boosting (XGB) Classifier_
        on Cross-Validation Set): ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(

```


```

    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Extreme Gradient
↳Boosting (XGB) Classifier on Cross-Validation Set): ROC-AUC and
↳Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Extreme Gradient
↳Boosting (XGB) Classifier on Cross-Validation Set): ROC-AUC and
↳Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```



Testing Set

```

[63]: with mlflow.start_run() as run:

    # Initializing the XGBoost classifier with default parameters
    xgb_model = XGBClassifier(use_label_encoder=False, eval_metric='logloss')

    # Training the model on the train dataset
    xgb_model.fit(X_train, y_train)

    # Predicting the output on the test dataset
    y_pred_final_xgb = xgb_model.predict(X_test)
    # Get probability scores for AUC calculation

```

```

y_scores_xgb = xgb_model.predict_proba(X_test)[: , 1]

# Calculate metrics
accuracy = accuracy_score(y_test, y_pred_final_xgb)
cm = confusion_matrix(y_test, y_pred_final_xgb)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

precision = precision_score(y_test, y_pred_final_xgb)
f1 = f1_score(y_test, y_pred_final_xgb)
roc_auc = roc_auc_score(y_test, y_scores_xgb) # Use probability scores for
↪AUC

# Update ensemble_df DataFrame
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "XGBoost Classifier",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the test set
print("XGBoost Classifier: ")
print(classification_report(y_test, y_pred_final_xgb))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(xgb_model, "XGBoost Classifier")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(xgb_model.get_params())
mlflow.log_metrics(
    {

```

```

        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)
# Save the model to MLflow
#shutil.rmtree("XGBoost Classifier", ignore_errors=True)
#mlflow.xgboost.save_model(xgb_model, "XGBoost Classifier")

signature = infer_signature(X_cv, y_pred_final_xgb)

# Log the sklearn model and register as version 1
mlflow.xgboost.log_model(
    xgb_model=xgb_model,
    artifact_path="xgboost-model",
    signature=signature,
    registered_model_name="xgboost-clf-model",
)

```

XGBoost Classifier:

	precision	recall	f1-score	support
0	0.92	0.86	0.89	28
1	0.85	0.92	0.88	24
accuracy			0.88	52
macro avg	0.88	0.89	0.88	52
weighted avg	0.89	0.88	0.88	52

Confusion Matrix:

```
[[24  4]
 [ 2 22]]
```

Run ID: c0c36305d5d04380a167c883eaa8ece8

Registered model 'xgboost-clf-model' already exists. Creating a new version of this model...

2024/04/13 12:37:14 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: xgboost-clf-model, version 3

Created version '3' of model 'xgboost-clf-model'.

```
[64]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_xgb)
roc_auc_val = auc(fpr, tpr)
```

```

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_xgb)
pr_auc = average_precision_score(y_test, y_scores_xgb)
f1 = f1_score(y_test, y_pred_final_xgb)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tp,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=recall,
        y=precision,
        mode="lines",
        name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=2,
)

```

```

fig.add_annotation(
    x=1.2,
    y=-0.15,
    xref="paper",
    yref="paper",
    text=f"Average Precision = {pr_auc:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
fig.add_annotation(
    x=1.2,
    y=-0.20,
    xref="paper",
    yref="paper",
    text=f"F1-Score = {f1:.2f}",
    showarrow=False,
    font=dict(size=15, color="blue"),
)

# Update layout
fig.update_layout(
    title_text="Model Performance (Extreme Gradient Boosting (XGB) Classifier):  

    ↳ ROC-AUC and Precision-Recall Curves",
    width=1200,
    height=600,
)
fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
fig.update_xaxes(title_text="Recall", row=1, col=2)
fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Extreme Gradient  

    ↳ Boosting (XGB) Classifier): ROC-AUC and Precision-Recall Curves.png") #png  

    ↳ format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Extreme Gradient  


    ↳ Boosting (XGB) Classifier): ROC-AUC and Precision-Recall Curves.svg") #svg  

    ↳ format

# Display the plots side-by-side
fig.show()

```


Model Performance (Extreme Gradient Boosting (XGB) Classifier): ROC-AUC and Precision-Recall Curves



Stacking Classifier (using Baseline Models)

Cross-Validation Set

```
[65]: with mlflow.start_run() as run:

    # Define baseline learners
    base_learners = [
        ("lr_clf_cv", lr_clf_cv),
        (
            "svm_model_cv",
            svm_model_cv,
        ), # Ensure svm_model is trained with probability=True for AUC scoring
        ("knn_model_cv", knn_model_cv),
        ("dt_clf_cv", dt_clf_cv),
    ]

    # Initialize the Stacking Classifier with the base learners and a final
    ↪ estimator
    final_estimator = DecisionTreeClassifier(random_state=42)
    stack_baseline_models_cv = StackingClassifier(
        estimators=base_learners, final_estimator=final_estimator, cv=5
    )

    # Fit the stack model
    stack_baseline_models_cv.fit(X_train, y_train)

    # Predict on the test set
```

```

y_pred_final_am_cv = stack_baseline_models_cv.predict(X_cv)
y_scores_am_cv = stack_baseline_models_cv.predict_proba(X_cv)[
    :, 1
] # Get probabilities for AUC scoring

# Calculate metrics
accuracy = accuracy_score(y_cv, y_pred_final_am_cv)
cm = confusion_matrix(y_cv, y_pred_final_am_cv)
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

roc_score = roc_auc_score(y_cv, y_scores_am_cv)
fpr, tpr, _ = roc_curve(y_cv, y_scores_am_cv)
roc_auc = auc(fpr, tpr)
precision = precision_score(y_cv, y_pred_final_am_cv) # Corrected to use
↳ predictions
f1 = f1_score(y_cv, y_pred_final_am_cv)

# Update DataFrame or similar storage
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Stacking Baseline Models CV",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_score,
        "F1-Score": f1,
        "Precision": precision,
    },
    ignore_index=True,
)

# Evaluate the model on the validation set
print("Stacking Baseline Models CV: ")
print(classification_report(y_cv, y_pred_final_am_cv))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow

```

```

mlflow.sklearn.log_model(stack_baseline_models_cv, "Stacking Baseline_
↳Models CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(stack_baseline_models_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Stacking Baseline Models CV", ignore_errors=True)
#mlflow.sklearn.save_model(stack_baseline_models_cv, "Stacking Baseline_
↳Models CV")

signature = infer_signature(X_cv, y_pred_final_am_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=stack_baseline_models_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-stack-baseline-cv-model",
)

```

Stacking Baseline Models CV:

	precision	recall	f1-score	support
0	0.63	0.76	0.69	25
1	0.74	0.61	0.67	28
accuracy			0.68	53
macro avg	0.69	0.68	0.68	53
weighted avg	0.69	0.68	0.68	53

Confusion Matrix:

```
[[19  6]
 [11 17]]
```

Run ID: 0a8eee265ed34a7fabd5e732d878028e

Registered model 'sk-learn-stack-baseline-cv-model' already exists. Creating a new version of this model...

2024/04/13 12:37:58 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-stack-baseline-cv-model, version 3

Created version '3' of model 'sk-learn-stack-baseline-cv-model'.

```
[66]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_am_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_am_cv)
pr_auc = average_precision_score(y_cv, y_scores_am_cv)
f1 = f1_score(y_cv, y_pred_final_am_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
    col=1,
)

# Add Precision-Recall Curve to the subplot
fig.add_trace(
```

```

        go.Scatter(
            x=recall,
            y=precision,
            mode="lines",
            name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
            fill="tozeroy",
        ),
        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Stacking Classifier {Baseline Models} on_
↪Cross-Validation Set): ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(
        margin=dict(b=100)
    )
    # Adjust the bottom margin to avoid cutting off annotations

    # Save plot


```

```

#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking
↳Classifier {Baseline Models} on Cross-Validation Set): ROC-AUC and
↳Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking
↳Classifier {Baseline Models} on Cross-Validation Set): ROC-AUC and
↳Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```



Testing Set

```

[67]: with mlflow.start_run() as run:

    # Define base learners
    baseline_learners = [
        ("lr_clf", lr_clf),
        ("svm_model", svm_model), # Ensure svm_model is trained with
↳probability=True for AUC scoring
        ("knn_model", knn_model),
        ("dt_clf", dt_clf),
    ]

    # Initialize the Stacking Classifier with the base learners and a final
↳estimator
    final_estimator = DecisionTreeClassifier(random_state=42)
    stack_baseline_models = StackingClassifier(

```

```

        estimators=baseline_learners, final_estimator=final_estimator, cv=5
    )

    # Fit the stack model
    stack_baseline_models.fit(X_train, y_train)

    # Predict on the test set
    y_pred_final_am = stack_baseline_models.predict(X_test)
    y_scores_am = stack_baseline_models.predict_proba(X_test)[
        :, 1
    ] # Get probabilities for AUC scoring

    # Calculate metrics
    accuracy = accuracy_score(y_test, y_pred_final_am)
    cm = confusion_matrix(y_test, y_pred_final_am)
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)

    roc_score = roc_auc_score(y_test, y_scores_am)
    fpr, tpr, _ = roc_curve(y_test, y_scores_am)
    roc_auc = auc(fpr, tpr)
    precision = precision_score(y_test, y_pred_final_am) # Corrected to use
    predictions
    f1 = f1_score(y_test, y_pred_final_am)

    # Update DataFrame or similar storage
    ensemble_df = ensemble_df._append(
        {
            "Ensemble Model": "Stacking Base Models",
            "Accuracy": accuracy,
            "Sensitivity": sensitivity,
            "Specificity": specificity,
            "False Positive Rate": FPR,
            "AUC Score": roc_score,
            "F1-Score": f1,
            "Precision": precision,
        },
        ignore_index=True,
    )

    # Evaluate the model on the test set

```

```

print("Stacking Baseline Models: ")
print(classification_report(y_test, y_pred_final_am))
print("")
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(stack_baseline_models, "Stacking Baseline Models")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(stack_baseline_models.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Stacking Baseline Models", ignore_errors=True)
#mlflow.sklearn.save_model(stack_baseline_models, "Stacking Baseline_
↳Models")

signature = infer_signature(X_cv, y_pred_final_am)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=stack_baseline_models,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-stack-baseline-model",
)

```

Stacking Baseline Models:

	precision	recall	f1-score	support
0	0.81	0.75	0.78	28
1	0.73	0.79	0.76	24
accuracy			0.77	52
macro avg	0.77	0.77	0.77	52
weighted avg	0.77	0.77	0.77	52

Confusion Matrix:

```
[[21  7]
 [ 5 19]]
```

Run ID: 7f8bbd0d8e7d4f6fa4f1896f96bca210

Registered model 'sk-learn-stack-baseline-model' already exists. Creating a new version of this model...

2024/04/13 12:38:37 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-stack-baseline-model, version 3

Created version '3' of model 'sk-learn-stack-baseline-model'.

```
[68]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_am)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_am)
pr_auc = average_precision_score(y_test, y_scores_am)
f1 = f1_score(y_test, y_pred_final_am)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
    row=1,
```

```

        col=1,
    )

    # Add Precision-Recall Curve to the subplot
    fig.add_trace(
        go.Scatter(
            x=recall,
            y=precision,
            mode="lines",
            name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
            fill="tozeroy",
        ),
        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Stacking Classifier {Baseline Models}):  

        ↳ ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(
        margin=dict(b=100)
    )

```

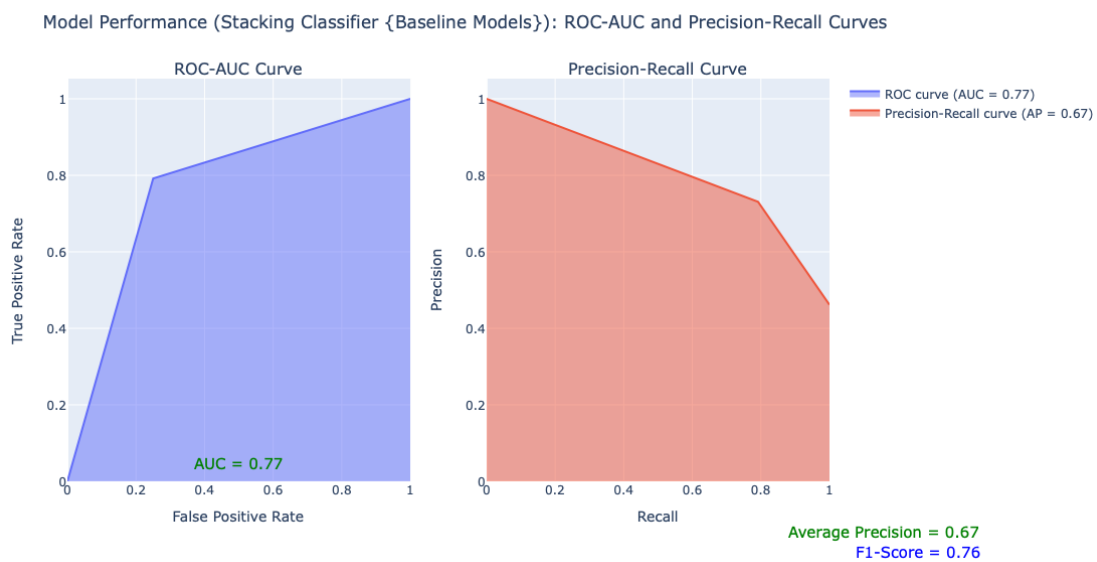
```

) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
↳Classifier {Baseline Models}): ROC-AUC and Precision-Recall Curves.png")
↳#png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
↳Classifier {Baseline Models}): ROC-AUC and Precision-Recall Curves.svg")
↳#svg format

# Display the plots side-by-side
fig.show()

```



Stacking Classifier (using ensemble models)

Cross-Validation Set

```
[69]: with mlflow.start_run() as run:
```

```

# Define the stacking ensemble model
stack_ensemble_models_cv = StackingClassifier(
    estimators=[
        ("xgb_model_cv", xgb_model_cv),
        ("gbm_model_cv", gbm_model_cv),
        ("vote_model_cv", vote_model_soft_cv),
        ("bag_model_cv", bag_model_cv),
    ],

```

```

        final_estimator=XGBClassifier(use_label_encoder=False,
eval_metric="logloss"),
        cv=5,
    )

    # Fit the stacking model
    stack_ensemble_models_cv.fit(X_train, y_train)

    # Predicting the final output using stacking
    y_pred_final_ams_cv = stack_ensemble_models_cv.predict(X_cv)

    # Get probability scores for AUC calculation
    y_scores_ams_cv = stack_ensemble_models_cv.predict_proba(X_cv)[: , 1]

    # Calculate accuracy score
    accuracy = accuracy_score(y_cv, y_pred_final_ams_cv)

    # Confusion Matrix
    cm = confusion_matrix(y_cv, y_pred_final_ams_cv)

    # Calculate specificity and other metrics
    TP = cm[1, 1]
    TN = cm[0, 0]
    FP = cm[0, 1]
    FN = cm[1, 0]
    sensitivity = TP / (TP + FN) # Recall
    specificity = TN / (TN + FP)
    # False Positive Rate (FPR)
    FPR = FP / (FP + TN)

    precision = precision_score(y_cv, y_pred_final_ams_cv)
    f1 = f1_score(y_cv, y_pred_final_ams_cv) # Calculate F1-Score
    roc_auc = roc_auc_score(y_cv, y_scores_ams_cv) # AUC score

    # Update ensemble_df DataFrame with new metrics
    ensemble_df = ensemble_df._append(
        {
            "Ensemble Model": "Stacking Ensemble Models CV",
            "Accuracy": accuracy,
            "Sensitivity": sensitivity,
            "Specificity": specificity,
            "False Positive Rate": FPR,
            "AUC Score": roc_auc,
            "F1-Score": f1,
            "Precision": precision,
        },
        ignore_index=True,

```

```

)

# Evaluate the model on the validation set
print("Stacking Ensemble Models CV: ")
print(classification_report(y_cv, y_pred_final_ams_cv))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(stack_ensemble_models_cv, "Stacking Ensemble_
↳Models CV")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(stack_ensemble_models_cv.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Stacking Ensemble Models CV", ignore_errors=True)
#mlflow.sklearn.save_model(stack_ensemble_models_cv, "Stacking Ensemble_
↳Models CV")

signature = infer_signature(X_cv, y_pred_final_ams_cv)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=stack_ensemble_models_cv,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-stack-ensemble-model-cv",
)

```

Stacking Ensemble Models CV:

	precision	recall	f1-score	support
0	0.82	0.72	0.77	25
1	0.77	0.86	0.81	28

accuracy			0.79	53
macro avg	0.80	0.79	0.79	53
weighted avg	0.79	0.79	0.79	53

Confusion Matrix:

```
[[18  7]
 [ 4 24]]
```

Run ID: c47d65ac669049688c5eefd6f79d13d2

Registered model 'sk-learn-stack-ensemble-model-cv' already exists. Creating a new version of this model...

2024/04/13 12:43:09 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-stack-ensemble-model-cv, version 2

Created version '2' of model 'sk-learn-stack-ensemble-model-cv'.

```
[70]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_cv, y_scores_ams_cv)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_cv, y_scores_ams_cv)
pr_auc = average_precision_score(y_cv, y_scores_ams_cv)
f1 = f1_score(y_cv, y_pred_final_ams_cv)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
```

```

        showarrow=False,
        font=dict(size=15, color="green"),
        row=1,
        col=1,
    )

    # Add Precision-Recall Curve to the subplot
    fig.add_trace(
        go.Scatter(
            x=recall,
            y=precision,
            mode="lines",
            name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
            fill="tozeroy",
        ),
        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Stacking Classifier {Ensemble Models} on   

        ↪ Cross-Validation): ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)

```


```

fig.update_yaxes(title_text="Precision", row=1, col=2)
fig.update_layout(
    margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
    ↳Classifier {Ensemble Models} on Cross-Validation): ROC-AUC and_
    ↳Precision-Recall Curves.png") #png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
    ↳Classifier {Ensemble Models} on Cross-Validation): ROC-AUC and_
    ↳Precision-Recall Curves.svg") #svg format

# Display the plots side-by-side
fig.show()

```



Testing Set

```

[71]: with mlflow.start_run() as run:

    # Define the stacking ensemble model
    stack_ensemble_models = StackingClassifier(
        estimators=[
            ("xgb_model", xgb_model),
            ("gbm_model", gbm_model),
            ("vote_model", vote_model_soft),
            ("bag_model", bag_model),

```



```

    ],
    final_estimator=XGBClassifier(use_label_encoder=False,
eval_metric="logloss"),
    cv=5,
)

# Fit the stacking model
stack_ensemble_models.fit(X_train, y_train)

# Predicting the final output using stacking
y_pred_final_ams = stack_ensemble_models.predict(X_test)

# Get probability scores for AUC calculation
y_scores_ams = stack_ensemble_models.predict_proba(X_test)[: , 1]

# Calculate accuracy score
accuracy = accuracy_score(y_test, y_pred_final_ams)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_final_ams)

# Calculate specificity and other metrics
TP = cm[1, 1]
TN = cm[0, 0]
FP = cm[0, 1]
FN = cm[1, 0]
sensitivity = TP / (TP + FN) # Recall
specificity = TN / (TN + FP)
# False Positive Rate (FPR)
FPR = FP / (FP + TN)

precision = precision_score(y_test, y_pred_final_ams)
f1 = f1_score(y_test, y_pred_final_ams) # Calculate F1-Score
roc_auc = roc_auc_score(y_test, y_scores_ams) # AUC score

# Update ensemble_df DataFrame with new metrics
ensemble_df = ensemble_df._append(
    {
        "Ensemble Model": "Stacking Ensemble Models",
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision,
    }, ignore_index=True

```

```

)

# Evaluate the model on the test set
print("Stacking Ensemble Models: ")
print(classification_report(y_test, y_pred_final_ams))
print("Confusion Matrix: ")
print(cm)

# Log the model parameters and metrics to MLflow
mlflow.sklearn.log_model(stack_ensemble_models, "Stacking Ensemble Models")
print("Run ID: {}".format(run.info.run_id))

mlflow.log_params(stack_ensemble_models.get_params())
mlflow.log_metrics(
    {
        "Accuracy": accuracy,
        "Sensitivity": sensitivity,
        "Specificity": specificity,
        "False Positive Rate": FPR,
        "AUC Score": roc_auc,
        "F1-Score": f1,
        "Precision": precision
    }
)

# Save the model to MLflow
#shutil.rmtree("Stacking Ensemble Models", ignore_errors=True)
#mlflow.sklearn.save_model(stack_ensemble_models, "Stacking Ensemble_
Models")

signature = infer_signature(X_cv, y_pred_final_ams)

# Log the sklearn model and register as version 1
mlflow.sklearn.log_model(
    sk_model=stack_ensemble_models,
    artifact_path="sklearn-model",
    signature=signature,
    registered_model_name="sk-learn-stack-ensemble-model",
)

```

Stacking Ensemble Models:

	precision	recall	f1-score	support
0	0.91	0.75	0.82	28
1	0.76	0.92	0.83	24
accuracy			0.83	52
macro avg	0.84	0.83	0.83	52

weighted avg 0.84 0.83 0.83 52

Confusion Matrix:

```
[[21  7]
 [ 2 22]]
```

Run ID: eb7357c89c7a455c9c00997c8d7ccc37

Registered model 'sk-learn-stack-ensemble-model' already exists. Creating a new version of this model...

2024/04/13 12:46:52 INFO mlflow.store.model_registry.abstract_store: Waiting up to 300 seconds for model version to finish creation. Model name: sk-learn-stack-ensemble-model, version 2

Created version '2' of model 'sk-learn-stack-ensemble-model'.

```
[72]: # Calculate metrics for ROC-AUC Curve
fpr, tpr, _ = roc_curve(y_test, y_scores_ams)
roc_auc_val = auc(fpr, tpr)

# Calculate metrics for Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y_test, y_scores_ams)
pr_auc = average_precision_score(y_test, y_scores_ams)
f1 = f1_score(y_test, y_pred_final_ams)

# Create subplots
fig = make_subplots(
    rows=1, cols=2, subplot_titles=("ROC-AUC Curve", "Precision-Recall Curve")
)

# Add ROC-AUC Curve to the subplot
fig.add_trace(
    go.Scatter(
        x=fpr,
        y=tpr,
        mode="lines",
        name=f"ROC curve (AUC = {roc_auc_val:.2f})",
        fill="tozeroy",
    ),
    row=1,
    col=1,
)
fig.add_annotation(
    x=0.5,
    y=0.05,
    xref="paper",
    yref="paper",
    text=f"AUC = {roc_auc_val:.2f}",
    showarrow=False,
    font=dict(size=15, color="green"),
)
```

```

        row=1,
        col=1,
    )

    # Add Precision-Recall Curve to the subplot
    fig.add_trace(
        go.Scatter(
            x=recall,
            y=precision,
            mode="lines",
            name=f"Precision-Recall curve (AP = {pr_auc:.2f})",
            fill="tozeroy",
        ),
        row=1,
        col=2,
    )
    fig.add_annotation(
        x=1.2,
        y=-0.15,
        xref="paper",
        yref="paper",
        text=f"Average Precision = {pr_auc:.2f}",
        showarrow=False,
        font=dict(size=15, color="green"),
    )
    fig.add_annotation(
        x=1.2,
        y=-0.20,
        xref="paper",
        yref="paper",
        text=f"F1-Score = {f1:.2f}",
        showarrow=False,
        font=dict(size=15, color="blue"),
    )

    # Update layout
    fig.update_layout(
        title_text="Model Performance (Stacking Classifier {Ensemble Models}):  

        ↳ ROC-AUC and Precision-Recall Curves",
        width=1200,
        height=600,
    )
    fig.update_xaxes(title_text="False Positive Rate", row=1, col=1)
    fig.update_yaxes(title_text="True Positive Rate", row=1, col=1)
    fig.update_xaxes(title_text="Recall", row=1, col=2)
    fig.update_yaxes(title_text="Precision", row=1, col=2)
    fig.update_layout(

```


```

margin=dict(b=100)
) # Adjust the bottom margin to avoid cutting off annotations

# Save plot
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
↪Classifier {Ensemble Models}): ROC-AUC and Precision-Recall Curves.png")_
↪#png format
#fig.write_image("./charts/Ensemble_Models/Model Performance (Stacking_
↪Classifier {Ensemble Models}): ROC-AUC and Precision-Recall Curves.svg")_
↪#svg format

# Display the plots side-by-side
fig.show()

```



```
[73]: ensemble_df.sort_values("Specificity", ascending=False)
```

```
[73]:
```

	Ensemble Model	Accuracy	Sensitivity	Specificity	False
	Positive Rate	Precision	AUC Score	F1-Score	
1	Voting Classifier (Soft)	0.903846	0.958333	0.857143	
0.142857		0.851852	0.885417	0.901961	
5	Gradient Boosting Classifier	0.884615	0.916667	0.857143	
0.142857		0.846154	0.961310	0.880000	
7	XGBoost Classifier	0.884615	0.916667	0.857143	
0.142857		0.846154	0.950893	0.880000	
4	Gradient Boosting Classifier CV	0.849057	0.892857	0.800000	
0.200000		0.833333	0.922857	0.862069	
6	XGBoost Classifier CV	0.830189	0.857143	0.800000	

0.200000	0.827586	0.917143	0.842105		
8	Stacking Baseline Models CV	0.679245		0.607143	0.760000
0.240000	0.739130	0.683571	0.666667		
0	Voting Classifier (Soft) CV	0.792453		0.821429	0.760000
0.240000	0.793103	0.880000	0.807018		
3	Bagging Classifier with SVM	0.826923		0.916667	0.750000
0.250000	0.758621	0.872024	0.830189		
11	Stacking Ensemble Models	0.826923		0.916667	0.750000
0.250000	0.758621	0.930060	0.830189		
9	Stacking Base Models	0.769231		0.791667	0.750000
0.250000	0.730769	0.770833	0.760000		
10	Stacking Ensemble Models CV	0.792453		0.857143	0.720000
0.280000	0.774194	0.900000	0.813559		
2	Bagging Classifier with SVM CV	0.773585		0.928571	0.600000
0.400000	0.722222	0.842857	0.812500		

Performance Metrics of Ensemble Models

```
[ ]: # Define a list of colors for the ensemble models
colors = [
    "rgba(255, 99, 132, 0.6)",
    "rgba(54, 162, 235, 0.6)",
    "rgba(255, 206, 86, 0.6)",
    "rgba(75, 192, 192, 0.6)",
    "rgba(153, 102, 255, 0.6)",
    "rgba(255, 159, 64, 0.6)",
]

# Map each ensemble model to a specific color
unique_models = ensemble_df["Ensemble Model"].unique()
color_map = {model: colors[i % len(colors)] for i, model in
    ↪enumerate(unique_models)}

# Create Subplots for the model performance using the 6 evaluation metrics
fig = make_subplots(
    rows=4,
    cols=2,
    subplot_titles=(
        "Accuracy vs. Ensemble Model",
        "Specificity vs. Ensemble Model",
        "Sensitivity vs. Ensemble Model",
        "False Positive Rate vs. Ensemble Model",
        "Precision vs. Ensemble Model",
        "F1-Score vs. Ensemble Model",
        "AUC Score vs. Ensemble Model",
    ),
    horizontal_spacing=0.15,
    vertical_spacing=0.15,
```

```

)

metrics = [
    "Accuracy",
    "Specificity",
    "Sensitivity",
    "False Positive Rate",
    "Precision",
    "F1-Score",
    "AUC Score",
]

plot_positions = [(1, 1), (1, 2), (2, 1), (2, 2), (3, 1), (3, 2), (4, 1)]

for metric, pos in zip(metrics, plot_positions):
    # Sort the DataFrame based on the current metric in descending order
    df_sorted = ensemble_df.sort_values(by=metric, ascending=False)

    # Extracting the sorted model names for consistent color mapping
    sorted_models = df_sorted["Ensemble Model"].unique()

    # Generate one bar for each model, now in sorted order
    for model in sorted_models:
        df_filtered = df_sorted[df_sorted["Ensemble Model"] == model]
        show_legend = (
            metric == "Accuracy"
        ) # Show legend only in the first subplot for clarity
        fig.add_trace(
            go.Bar(
                x=[model],
                y=df_filtered[metric],
                name=model,
                marker_color=color_map[model],
                text=df_filtered[metric].round(2),
                textposition="outside",
                showlegend=show_legend,
            ),
            row=pos[0],
            col=pos[1],
        )

# Update layout
fig.update_layout(
    height=1500,
    width=1100,
    title_text="Performance Metrics of Ensemble Models",
    showlegend=True,
    legend=dict(orientation="v", x=1.05, y=0.5),

```

```

        font=dict(size=10),
    )
fig.update_xaxes(tickangle=45)
fig.update_yaxes(range=[0, 1])

# Save plot
#fig.write_image("./charts/Ensemble_Models/Performance Metrics of Ensemble_
↳Models.png") #png format
#fig.write_image("./charts/Ensemble_Models/Performance Metrics of Ensemble_
↳Models.svg") #svg format

# Display the plot
fig.show()

```

The ***Extreme Gradient Boosting algorithm*** provided the best performance among the six ensemble models in terms of 5 (out of 6) evaluation metrics used.

Conclusively, it was found that the ensemble models outperformed the baseline models; while **Logistic Regression** being a base model had the highest Specificity score (0.85), the **Extreme Gradient Boosting algorithm (XGB)** being an ensemble model had the highest overall Specificity score of 0.89.